

UNIVERSITY OF HERTFORDSHIRE
SCHOOL OF COMPUTER SCIENCE

MODULAR BSc HONOURS IN COMPUTER SCIENCE

6COM0282 COMPUTER SCIENCE PROJECT

FINAL REPORT
APRIL 2013

WORK-STEALING SCHEDULING TECHNIQUES
APPLIED TO COMPUTING THE MANDELBROT SET

M.J.HAWES

SUPERVISED BY: COLIN EGAN

Abstract

Load-balancing of parallel computer programs is a key technique for improving performance. It also serves to displace some of the burden of fairly distributing work-load from the programmer. Work-stealing is a technique used to implicitly balance work-load across multiple processors throughout run-time. In this report the technique is explored and applied to an algorithm to compute a raster plane of the Mandelbrot set.

A Randomised Work-Stealing Algorithm which utilises a non-blocking double ended queue is implemented and successfully applied to a variation of the level-set approach to computing the Mandelbrot Algorithm.

The investigation also encompasses some exploration into fractals and the Mandelbrot Set as a means of defining part of the problem domain.

The report concludes with a number of findings which suggest that the work-stealing approach offers significant benefits such as; a performance increase versus an un-balanced parallel approach, good scalability as problem size increases, and a fair approach to choosing where work to be migrated is sourced. In addition a number of weaknesses of the approach are identified and discussed.

Acknowledgements

The author would like to acknowledge Colin Egan for his invaluable guidance throughout this project, and for his sound technical and academic advice.

Contents

1	Introduction	2
2	Background Research	5
2.1	Run-Time Scheduling Techniques for Multi-Threaded Computations . .	5
2.2	The Mandelbrot Set	5
2.3	The Work-Stealing Technique - Described in Depth	8
2.3.1	Blumofe and Leiserson - A Randomized Work-Stealing Algorithm	8
2.3.2	McGuiness - Render-Thread Algorithm	9
2.4	Tools	10
3	The Development and Implementation	13
3.1	General Design and Practices	13
3.2	An Algorithm to Compute the Mandelbrot Set	15
3.2.1	The Mandelbrot Module	15
3.3	A Randomised Work-Stealing Algorithm	17
3.3.1	The Deque Implementation	17
3.3.2	The Work-Stealing Mechanism	19
3.4	Additional Schemes	22
3.5	Output	25
3.6	Run-Time Tracing	26
3.7	User Interface	27
4	Discussion and Evaluation	29
4.1	Analysis of the Implemented Algorithms	29
4.1.1	Performance Analysis	29
4.1.2	Run-Time Analysis of the Randomised Work-Stealing Algorithm	31
4.2	Reflection	34
4.3	Further Work	36
4.4	Conclusion	38
5	Resources	39
6	Appendices	42
A	Glossary of Terms	43
B	Trace Output	46
C	Source Code	48

Chapter 1

Introduction

Multi-processor computers are now commonplace for both consumers and scientists alike, in replacement of uni-processors. This shift in architectural design demands a drastic shift in how programs for such platforms are constructed, from the standpoint of both the programmer and the language designer. Designing effective programs to utilise such hardware, using well established programming tools with sequential roots, is often difficult and riddled with pitfalls. Popular languages such as C, Java, and C++ require a good knowledge of multi-threaded algorithms, as well as a deeper understanding of the problem domain to produce a highly effective design. The programmer may be forced to sacrifice potential performance gain for a more manageable level of design complexity, or worse still; avoid writing parallel code all-together in the interest of robustness.

Techniques to implicitly optimise run-time performance, with a minimum impact on complexity of parallel program design, are useful for more easily producing robust implementations. One aspect which allows for such optimisation is run-time work-load balancing.

This project explores a branch of load balancing techniques known as work-stealing. This is applied to an algorithm which computes an approximation of the infinite mathematical set known as the Mandelbrot set.

Motivation

The author has a strong interest in high performance computing and tools which provide powerful, abstract interfaces with underlying multi-processor architecture. Previous work on the Single Assignment C [9] and S-Net [8] programming languages has served to inspire, and further the desire for knowledge of the subject area.

This project is motivated by the aspiration to better understand how modern computer architectures can be optimally utilised, and communicate these findings in a concise, comprehensive report.

In addition the author wishes to gain valuable experience of the project development and management process.

Aim

To investigate, through research and practical implementation, the effectiveness of work-stealing on a parallel algorithm which computes an approximation of the Mandelbrot set.

Objectives

This section outlines the deliverable items that are presented in this report. Core objectives are primary and more vital, advanced objectives are additional items.

Core Objectives

1. **Background Research**
2. **A Sequential Mandelbrot Set Algorithm**
3. **A Naïve Parallel Mandelbrot Set Algorithm**
4. **A Random Work-Stealing Mandelbrot Set Algorithm**
5. **Analysis of the Implemented Algorithms**

Advanced Objectives

6. Graphical Output
7. A Render Thread Work-Stealing Mandelbrot Set Algorithm
8. Work-Stealing Trace System

Achievements

The following table shows each objective and its completion status at the time this report was produced.

Objective	Status
1	Complete
2	Complete
3	Complete
4	Complete
5	Complete
6	Complete
7	Partially Complete
8	Complete

Table 1.1: Table of Achievements

Report Structure

There are five remaining chapters in this report. In chapter two a detailed review of the problem background is presented. Chapter three details the practical implementation, which this report is based on, and the manner in which it was carried out. Chapter four offers an evaluation of the project as a whole. This includes analysis of the implemented software and a comparison of the implemented schemes. Chapter five lists the resources referenced in this report.

In chapter six the appendices are presented. A glossary of terms is presented as Appendix A. This consists of a list describing major concepts surrounding the field of study, and directs the reader to their page of first occurrence. Words emphasised in bold typeface indicate the first appearance of a glossary term. Some example trace

output is listed in Appendix B. The source code for all software developed to fulfil the objectives described above is presented as Appendix C.

Chapter 2

Background Research

2.1 Run-Time Scheduling Techniques for Multi-Threaded Computations

This section briefly describes the problems associated with **scheduling multi-threaded algorithms** at run-time and the major paradigms that have surfaced.

To efficiently utilise a parallel computer architecture it is desirable to minimise the amount of time a processor spends idle or performing other logistical tasks, i.e not doing work. When a computation's concurrent sub-tasks (or **threads**) incur a regular cost in processor time, each processor can simply have the same amount of work assigned to them. When the computation has more irregular or dynamically growing sub-tasks, a problem arises resulting in processors becoming idle while others still remain working. The solution to this problem is referred to as **load-balancing** and can be described as a form of dynamic scheduling, that ensures each processor spends approximately the same amount of time working. This means processors generally spend less time idle, however have to deal with scheduling overheads as a trade-off.

When considering the scheduling of multi-threaded computations, two major load balancing techniques have been used. These are **work-sharing** and **work-stealing**.

- **Work-Sharing:** A processor which creates new work attempts to migrate it to another underutilised processor at creation time.
- **Work-Stealing:** A processor which is starved of work attempts to “steal” work from other processors.

Both techniques intend to promote balanced work-load across all processors, however in Work-Stealing the frequency of work migrations is lower. When all processors have a high work-load and no need to “steal” this becomes useful because threads need not get migrated at all. With work-sharing work migration occurs each time new work is created [13]. This also suggests that work-stealing promotes better **locality** and grouping of sub-tasks, as spawned work stays with the same processor until stolen.

2.2 The Mandelbrot Set

The Mandelbrot set is a set of complex numbers which when plotted produce a spectacular and recognisable shape as illustrated in figure 2.1. It is often presented as a colourful and striking image and has been described by some as the most beautiful object in all of mathematics [19, p. 234]. The **complex-numbers** that comprise the set are closely related to julia-sets. In-fact the Mandelbrot set can be described as a

catalogue of Julia Sets which, when plotted, all points are connected forming a single unbroken shape [12, p. 177].

The set is named for the mathematician Benoit Mandelbrot, who discovered it in 1980 [12, 25]. He was a pioneer in the study of fractal geometry and also coined the term **fractal**, of which both the Mandelbrot set, and Julia sets are examples of.

In this section I will give a more detailed explanation of the areas mentioned here.




Figure 2.1: A rendering of The Mandelbrot Set generated using the program “fraqtive” [2].

Fractals and Self-Similarity

A fractal is a means of describing shapes which are more complex than a Euclidean shape. The leaves of a pine tree or the forks of a lightning bolt are obvious examples of real things that fractals allow us to more faithfully describe. These **real-world-fractals** are similar to **mathematical-fractals**, of which the Mandelbrot set is an example, but differ in that they do not display the property of **scale-invariance**.

Fractals have a fractional dimension. Unlike shapes with topological dimension, for instance a two dimensional square, a fractals dimension is of a non integer value.

A property of fractals (but not all) is **self-similarity**, where the shape is comprised of smaller “copies” of itself. This is known as **exact self-similarity** and means the shape is identical at any scale. A well-known example of this is the Triadic Koch Snowflake which is a fractal constructed using equilateral triangles. It is important to note here that The Mandelbrot set does not quite show the same property, it is said to be **quasi self-similar**. This means the shape is approximately similar at all scales, in that the shape is replicated but in a slightly distorted form with each “copy”.

It turns out there are many rather useful applications for fractals. To name a few; computer graphics used in video games and film (notably Star Trek II: The Wrath of Khan [12, p. 8]), military hardware design [17], and measuring the length of a coastline [23].

Julia Sets

To understand the basis of the Mandelbrot set it is first necessary to understand it's relation to julia-sets. The function in equation 2.1 is iterated infinitely where c is fixed. The filled julia-set is comprised of all values of z_0 where the result is bounded and does not tend towards infinity. The julia-set is comprised of those members of the filled julia-set which lie on the boundary [19]. In the interest of keeping this report readable, and because filled Julia Sets are more relevant, filled julia-sets will be referred to simply as julia-sets.

$$f(z) = z^2 + c \quad (2.1)$$

The Mandelbrot Set is related to julia-sets in which the values c and z used are expressed as a complex-number. Figure 2.2 illustrates some examples of such sets.




Figure 2.2: Two Julia Sets rendered using the program “fraqtive” [2]. Figure 2.2i is a member of the Mandelbrot set, figure 2.2ii is not.

The Mandelbrot Set and How it is Computed

The set is comprised of those julia-sets which are connected. In order to determine whether a Julia Set possesses this property, we need only compute the result for z_0 . If this tends towards infinity the value c is not a member of the Mandelbrot Set. If the result is bounded, then it is a member. This is known as the **critical-orbit** and is useful because it means we do not have to compute the entire Julia Set for each value of c .

So the Mandelbrot set can be computed by iterating all possible values of c for the function in equation 2.1 where z is the critical-orbit. Because the set of all possible values of c is infinite, and computers have a finite amount of resources, this set needs to be approximated. This can be done using a raster plane which takes samples of the complex plane at regular intervals.

There are many algorithms including the Level set method [12, p. 188] and Continuous Potential Method [12, p. 191]. The former uses a raster plane to produce an approximation as described above, and the latter produces a smooth surface representation.

Listing 2.1 describes a variation of the level set method in pseudo code. The algorithm is simple and provides a gradient effect of equipotential curves on the output image which is visually pleasing.

```
1 compute_mandelbrot()
2   FOR y = 0 TO height - 1 DO
3     c_im := im_min + y * (im_max - im_min) / (height - 1)
4     FOR x = 0 TO width - 1 DO
5       c_re := re_min + x * (re_max - re_min) / (width - 1)
6       output_plane[x][y] := level_set(c_re, c_im, max_iterations)
7     END FOR
8   END FOR
9 END
10
11 level_set(c_re, c_im, max_iterations)
12   z_re := c_re
13   z_im := c_im
14
15   FOR i = 0 TO max_iterations DO
16     IF( z_re^2 + z_im^2 > 4) THEN
17       RETURN i
18     END IF
19
20     tmp_im = 2 * z_re * z_im + c_im
21     z_re := z_re^2 - z_im^2 + c_re
22     z_im := tmp_im
23
24   END FOR
25   RETURN max_iterations
26 END
```

Listing 2.1: A sequential algorithm to compute the Mandelbrot Set

2.3 The Work-Stealing Technique - Described in Depth

As described above, work-stealing is a load-balancing technique which allows work starved processors to acquire scheduled work from other processors.

Each processor has a number of assigned tasks to complete. In general, a processor acquires its work from here. However, once these tasks are exhausted, the processor becomes a **thief** and a **victim** is chosen to steal from. The method used to choose a victim is implementation specific, for instance some implementations adopt a random scheme [10, 13, 22]. If the processor successfully steals work it relinquishes its thief state and returns to doing work. If the steal attempt is unsuccessful, for instance when the victim has no work or is blocked, the processor tries again until it is determined that there is no work remaining in the entire network.

Figure 2.3 illustrates the result of a successful steal operation in which work-starved processor *p1* transfers a piece of work from *p0*'s work list to its own.

Research has been conducted to explore its application in programming languages such as Cilk [13], parallel programming libraries such as Hood [14] and Factory [30], and large scale heterogeneous systems such as computational clouds [20].

This section explores some schemes used to implement work-stealing in various settings. It is focused on overall design and techniques presented in related literature.

2.3.1 Blumofe and Leiserson - A Randomized Work-Stealing Algorithm

This scheme is geared towards computation of dynamically growing, fully strict, multi-threaded computations and is applied to the CILK programming language and its runtime system [13].

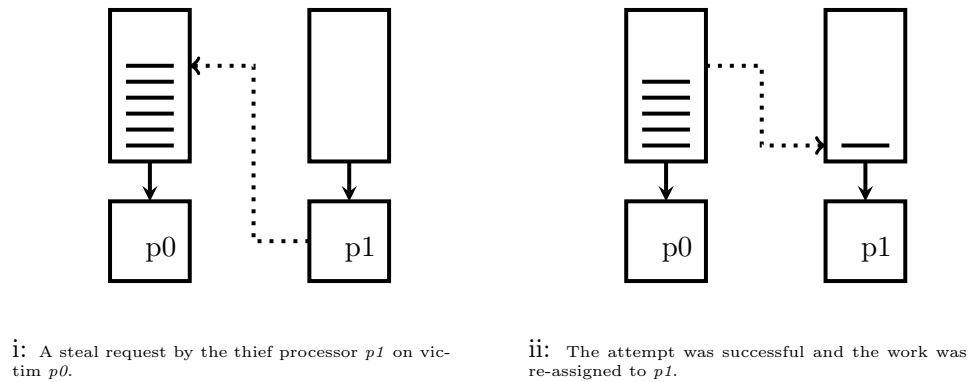


Figure 2.3: A successful steal operation between a thief and its victim.

Each thread maintains a **ready-deque**; a double-ended queue of work waiting to be processed. Accesses to this queue are made either at the top (for a steal operation), or at the bottom (for a push operation or when the next piece of work is required). A thread becomes a thief when its ready-deque is empty and randomly selects a victim; a thread to attempt to steal work from. If this **steal-operation** is successful it pushes the stolen work onto the bottom of its ready-deque and becomes a worker again. If not it tries, at random, to find another victim.

The ready-deque can be implemented in such a way that a thread need not be stopped in order for a steal operation to occur. This property is known as **non-blocking** and only requires that the top end of the deque has atomic access, while the bottom can freely be accessed by the thread which owns the deque [10]. This is useful because it reduces the overheads of a steal operation in that a working thread generally does not get interrupted. Further still, a non-blocking deque can be made more efficient through use of a **circular array** [15].

Because of the setting this scheme is designed with in mind, the algorithm needs to consider that a piece of work can spawn children dynamically, which it may depend on completing to continue. When computing the Mandelbrot set in a concurrent environment, no such consideration is required as each point can be independently processed.

2.3.2 McGuiness - Render-Thread Algorithm

This scheme is presented as part of McGuiness' Masters Thesis [26], and is suited for parallel computation where each unit of work is independent from any other and can be indexed in a list. Computation of the Mandelbrot set is given as an application of the algorithm.

The algorithm uses a set of **worker-threads** (referred to by McGuiness as render-threads), and a single **monitor-thread**, to control the distribution of work. Each worker-thread is initially given an equal share of the overall work-load before starting.

Each worker-thread maintains an estimated completion time for its assigned work-load. This is initially set to the maximum possible value and is iteratively refined by calculating the average time taken to complete a piece of work. This metric is used as a policy for deciding which thread is the most suitable candidate for the victim of a work-stealing operation.

When a worker-thread completes its assigned work a work-completed signal is generated, its estimated completion time is set to 0, and the thread is stopped. This thread will be referred to as the thief. When the monitor thread detects this signal, it searches for the worker-thread with the longest estimated completion time, which will be referred to as the victim. The monitor-thread waits for the victim to complete its

current piece of work before stopping it. Its workload is then halved, having the other half re-assigned to the thief. Both the victim and the thief are restarted and continue doing work. The monitor-thread returns to waiting for another work-completed signal and the process is repeated until no work remains.

Discussion of the Presented Schemes

The key difference between the Render-Thread Algorithm and the Randomized Work-Stealing Algorithm is that a worker-thread does not maintain a queue of work, but simply has a range of indices assigned. This all but eliminates the overheads associated with initialising and maintaining a potentially costly data-structure, but makes a non-blocking implementation difficult. It also requires a monitor-thread to manage work migration which reduces the maximum number of threads performing work by one. Another limiting factor is that only one thread may perform work-stealing at any given time. The Randomized Work-Stealing Algorithm has no such limitation.

2.4 Tools

This section discusses some of the programming and general tools which were considered for use in the implementation and for the good of the project as a whole. The tools discussed in this section are all viable options for a Linux platform.

Programming Languages

- **C:** A general purpose imperative language which is very widely used. It is statically typed but weakly enforces type errors and allows low level access to memory, which makes the programmer responsible for reclaiming used memory. This allows the programmer a high level of control but reduces type safety. It is not object oriented and lacks many features of more modern high level language features.

C is well suited for implementing efficient, high performance programs due to its relatively low level of abstraction compared to more modern languages. It is often associated with system programming because of this. A number of parallel programming libraries are available amongst other useful libraries.

GCC is an open source, industry standard compiler and is freely available. It is well documented and supported.

- **C++:** A general purpose multi-paradigm language inspired by C. C++ adds object oriented programming (amongst other features) to the majority of C's syntax and features. This makes it more suitable for implementing larger application systems with the same benefits of low-level memory manipulation that C boasts.

Some useful features include inheritance of types, dynamic polymorphism, and templates (including an extensive library of useful template classes). A number of parallel programming libraries are available and are similar if not the same as the equivalent C versions.

G++ is a variant of the GCC compiler and shares a number of its properties.

- **Java:** Primarily an object oriented language which borrows features from other paradigms. It is highly portable because it targets java bytecode which runs on the java virtual machine, available on most platforms. It is heavily influenced by C and C++ but offers few low-level facilities available in said languages. For

instance pointers are not available and a garbage collector is used to move the responsibility of memory management from the programmer, to the run-time system. Although this is a useful tool, the consequence is a considerable run-time overhead.

Java is more suited to application programming than the likes of C and C++. Its extensive API offers concurrent programming classes.

An open source compiler suite is available as well as a propriety version. These are both widely used and well documented.

Parallel Programming Libraries

In order to implement Work-Stealing, support for programming threads with a suitable level of control is required.

- **POSIX Threads (pthreads):**

Provides low level manipulation of threads for the C programming language [11]. It is a library based on IEEE standard 1003.1. Thread programming is achieved by specifying a function in which to run in parallel. Thread synchronisation is supported through use of a set of functions and data structures provided; such as **mutexes**, **barriers**, and **condition-variables**.

- **Open Multiprocessing (OpenMP):**

Provides abstract thread programming interfaces for C, C++, and fortran. In general OpenMP only allows coarse grained manipulation of threads through features such as parallel loops [21].

- **Java Threads:**

An object oriented approach to multi-threaded programming. Java provides a Thread class which the programmer can specialise to perform the desired task. Synchronisation is generally done using implicit locks through use of the synchronised modifier.

Graphical Output

For the purpose of demonstrating that the program correctly generates a raster plane of the Mandelbrot set, a graphical representation of the plane is output.

- **PPM Output File:**

The simplest option is to output to a Portable Pixel Map (PPM) file. It is text-file based and easy to implement but produces rather large files. The process of outputting to a text file is inherently serial in nature, so with large image resolutions processing takes a long time. There is support for both grey-scale (Portable Grey-scale Map format) and colour images. The advantage of using this method is the portability. No libraries or extras are required and most image viewers will read the file. [7]

- **GNU Plot:**

A graph plotting package available for multiple operating systems. Supports screen display or file output of both 2d and 3d graphics [5]. There are programming interfaces available for various languages such as C [18], C++ [6], and Java [1]. GNU Plot needs to be installed on the machine in order to use it.

- **OpenGL - glut:**

Glut is a framework for providing simple, cross-platform, GUI window control in conjunction with OpenGL. Bindings are available for various languages, including C and C++ [4]. A free implementation called “freeglut” is available and can be installed on linux [3]. This method requires OpenGL and an implementation of Glut be available on the machine that the program is run on.

Chapter 3

The Development and Implementation

3.1 General Design and Practices

Development Process Model

A feature driven development (FDD) style approach to development is adopted for this project. FDD is a branch of project management derived from the school of agile process models [27]. It is almost ideal for this project due to the small scale, short product life-cycle, and the flexibility needed when requirements change as the project transpires.

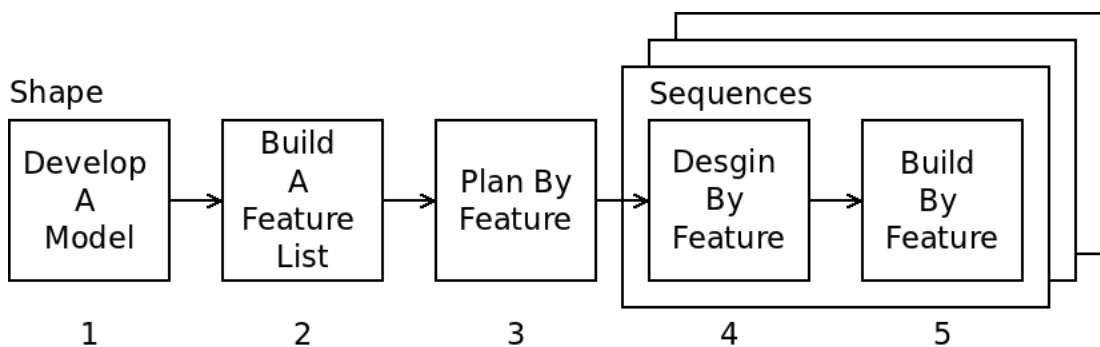


Figure 3.1: The five stages of the feature driven development cycle. The overlapping rectangles behind stages four and five indicate sequences that can be implemented concurrently. The diagram is derived from material presented here [16, p. 90].

Some of the notable features of FDD which make it attractive for such a project are described as follows:

- Short, iterative cycles of development.
- Working results delivered by each iteration.
- Flexibility for developers and clients alike, allowing for easy adaptation to changing requirements.
- Product quality is emphasised at each step.
- Development stages may overlap, i.e. individual features can be produced concurrently.

Major Tool Choices

The implementation is realised using the C programming language, the pthreads library, and PPM files for graphical output. These choices are based on the tools examined in section 2.4.

The C programming language is chosen for its relatively low-level of abstraction and level of control over memory management¹. It is preferable over C++ and Java because features such as object oriented programming are deemed unnecessary run-time overheads. The implementation is relatively small-scale, so C will suffice in terms of managing the code complexity.

The pthreads library is selected for its low-level control compared to the likes of OpenMP and Java Threads. The implementation of the Randomised Work-Stealing algorithm only requires control of simple thread synchronisation (i.e. mutexes) and nothing more.

The text based PPM file format is chosen for its simplicity and portability. Producing a graphical representation of the Mandelbrot set is not the primary purpose of this project, but is still a desirable feature for verification and demonstration purposes.

Other Tools Used

Listed here are the major utility tools which have been used to make this project of a better over-all quality.

- **L^AT_EX:**
A document preparation mark-up language which produces professional and consistent documents.
- **GNU Make:**
A tool to aid quick and easy building of a project. It uses a series of rules and dependencies to determine the order in which a project can be built, which are described in an accompanying Makefile.

This is useful for producing an executable from source code, as well as compiling a L^AT_EX document in conjunction with bibtex to produce a pdf.
- **Git:**
A distributed version control system which provides version tracking capabilities. It has the benefit of providing a means of backing up, as-well as maintaining synchronisation of project files across multiple machines. It also serves to document the progress of a project through commit messages.

Coding Conventions

All presented code follows strict conventions to ensure readability. Following is a list of conventions used.

- **Names:**
All variable and function names are given in lower-case with multiple words separated by an underscore. Function names declared in a module are prefixed with an associated acronym. For instance, functions declared in 'deque.c' start with 'de_'. Type definitions are named in the same fashion as a variable, followed by a trailing '_t', to match the style used by the pthreads library.

¹This decision is contrary to the choice (i.e. C++) documented in the project proposal.

- **Parentheses:**

For functions and iteration statements the opening parenthesis is placed on a new line, in any other case it can be placed on the same line. The closing parenthesis is always placed on its own line.

- **Constants and Pre-Processor Directives:**

Pre-processor directives are declared in the accompanying header file. All names are in upper-case with words separated by an underscore.

- **Commenting**

Comments are used to describe code sections which are challenging to understand or where complicated constructs are used. Where code is simple enough to be self documenting comments are used sparingly. In general the block style is used (i.e. `/* ... */`).

Modular Design

The implementation takes the form of a modular design. The module which has the functionality to compute the Mandelbrot set has an interface to which a separate ‘scheduling’ module can be attached. The Makefile holds a rule for each scheduling module and builds a binary for each. This approach provides several benefits.

- It ensures all scheduling modules use exactly the same scheme to compute the Mandelbrot set, making them more comparable.
- It allows for a common user-interface for each executable, independent of the scheduler back-end.
- It allows for alteration to the mandelbrot module to be made with ease.

3.2 An Algorithm to Compute the Mandelbrot Set

This section describes the Mandelbrot module, which provides an interface for run-time scheduler implementations to compute a raster-plane of the Mandelbrot set line by line. This interface is provided primarily by the `compute_line` function. The module requires that any attached scheduler module implements the functions `ws_initialise_threads` and `ws_start_threads`. The former should be used for any set-up (e.g initialising global variables), and the latter is used to start processing the raster plane.

3.2.1 The Mandelbrot Module

The key functions used in the mandelbrot module are described as follows:

- **compute_line:**

Iterates a single line of the x axis on the raster plane, specified by a parameter y , and assigns the result to the corresponding members of the plane matrix. It also accepts a ‘`t_id`’ (thread id) parameter, which is used for visualising regions of work completed per thread. This is used as a point of interface for the attached scheduler module, which is responsible for iteration of the y axis. The implication of this design means that the minimum granularity of a work item is one ‘line’.

The function converts the co-ordinates x and y into a corresponding complex number c using the `convert_x_coord` and `convert_y_coord` functions respectively. Should the value of c lie outside the radius of two from the origin, the point is assigned the constant `PPM_BLACK`. Its thread id is assigned `WORKER_COUNT`

to indicate to indicate that the graphical output should ignore any thread information for this pixel. Otherwise the pixel is assigned the value returned from the `is_member` function and the thread id is assigned the value of the `t_id` parameter.

- **is_member:**

Returns a value between 0 and the constant `MAX_ITERATIONS`. It accepts a complex number c which corresponds to one point on the complex plane and one pixel of the raster plane. This function determines (approximately) whether point c tends towards infinity.

The variable z is initially assigned the value of c . The function iterates until either the `MAX_ITERATION` count is reached or $\sqrt{z_r^2 + z_i^2} > 2$, which is simplified to $z_r * z_r + z_i * z_i > 4$ to avoid the 'sqrt' function for efficiency reasons. The former condition indicates that c is a member of the set and a constant to indicate this is returned. The latter indicates that c is not a member of the set and the number of iterations is returned, which produces a gradient effect on the output image. At the end of each iteration the next value of z is obtained by calling the `julia_func` function, passing the current value of z and the value of c .

- **julia_func:**

Computes the function shown in equation 2.1. It accepts two complex numbers; z and c .

The returned complex number is calculated by the following two expressions; for the imaginary part $res_i = 2 * z_r * z_i + c_i$ and for the real part $res_r = (z_r * z_r) - (z_i * z_i) + c_r$.

- **convert_y_coord** and **convert_x_coord:**

Collectively converts a coordinate of the raster plane into its corresponding point on the complex plane. The `convert_y_coord` returns the imaginary part and the `convert_x_coord` returns the real part. Both return a double type. The two are separated to avoid unnecessary conversion of the y coordinate in the `compute_line` function, as this only need occur once per 'line'.

The algorithm is based on the level set method [12, p. 188]. It is optimised slightly by checking that the point is not outside the radius of two before calling the `is_member` function. This is because no member of the set lies outside this radius thus saving us iterating for such a point unnecessarily.

Each pixel of the raster plane maps to a point on the complex plane. Figure 3.2 demonstrates this.

Several parameters are used by the algorithm to construct the raster plane.

- **HEIGHT** and **WIDTH:**

Constants which define the dimensions of the raster plane i.e the maximum values of y and x respectively.

- **MAX_ITERATIONS:**

A constant iteration limit which is used to determine if a point lies within the set. This is used by the `is_member` function.

- **The c_max and c_min:**

Complex numbers which define the limits of the complex plane of which the raster plane samples.

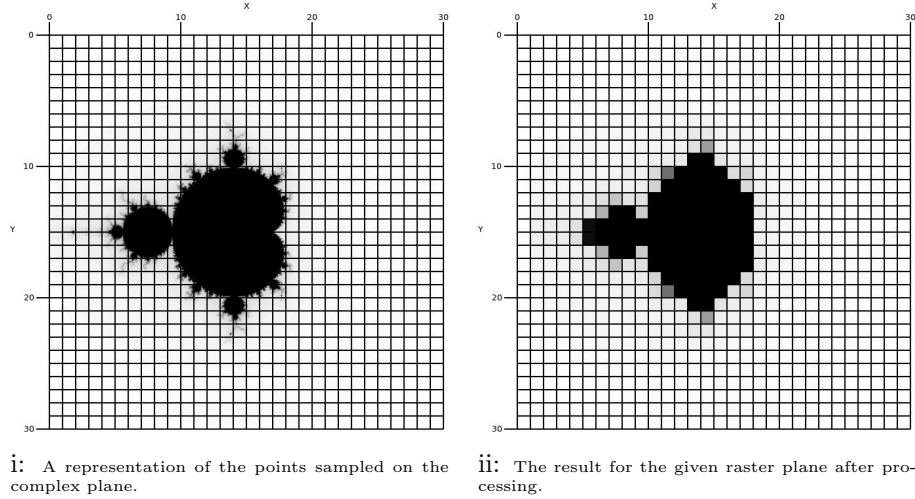


Figure 3.2: A visual representation of an example raster plane where HEIGHT and WIDTH equal thirty and MAX_ITERATIONS equals seventy. Figure 3.2ii shows how the raster plane produces an approximation of the complex plane.

- **The *c_factor*:**

A complex number which determines the difference between a sampled point of the complex plane, and an adjacent sample. It is calculated using the dimensions of the raster plane, and the minimum and maximum values of *c*.

$$cfactor_r = cmax_r - cmin_r / rasterwidth - 1$$

$$cfactor_i = cmax_i - cmin_i / rasterheight - 1$$

3.3 A Randomised Work-Stealing Algorithm

This section describes the implementation of the Randomised Work-Stealing algorithm based on the scheme described in section 2.3.1. This implementation uses four threads, each maintaining its own ready-deque. The implemented ready-deque is based on the scheme presented in [15].

3.3.1 The Deque Implementation

Three key operations are made on a ready-deque. These are `de_pop_bottom`, `de_push_bottom`, and `de_steal`.

- **`de_pop_bottom`:**

Accepts a Deque pointer and returns a Line from the front of the deque, or a Line signalling an empty Deque.

The bottom counter of the Deque passed to the function is decremented regardless of the outcome. If the deque is empty a Line with a status of `LINE_EMPTY` is returned, expecting the client code to handle this appropriately. If the Deque has more than one item, it may be shrunk and the Line indexed using the bottom counter is returned. If the Deque only holds one item and the `top_mutex` is locked, then a simultaneous steal operation has claimed the last item. In this case a Line with a status of `LINE_EMPTY` is returned. Otherwise the remaining bottom item is returned.

- **de_steal:**

Accepts a Deque pointer and attempts to return a Line from the back of the deque. Otherwise a Line with status `LINE_EMPTY` or `LINE_ABORT` is returned.

If the Deque has only one item remaining a Line with the status `LINE_EMPTY` is returned. This is tested before the `top_mutex` is locked in order to avoid unnecessary blocking of the `pop_bottom` operation. In any other case an attempt to lock the `top_mutex` is made. If this fails then a simultaneous `pop_bottom` operation has locked the deque item, and a Line with the status `LINE_ABORT` is returned. Otherwise the item at the top of the Deque is returned and the top counter is incremented.

- **de_push_bottom:**

Accepts a Deque pointer and a Line to push onto the bottom of the queue.

It attempts to re-size the array (if it needs to) and increments the bottom counter whilst placing the Line passed to it on the bottom of the deque.

Both the `de_pop_bottom` and `de_steal` operations, in some situations, need to ensure that the thread has exclusive access of the top index of the deque. This is achieved using the `pthread_mutex_trylock` function, which accepts a mutex and returns 0 should lock be successful. Otherwise, for instance when the mutex is locked by another thread, an error value is returned. This allows the thread to test the state of the mutex, and continue execution without blocking (waiting for the mutex to become un-locked).

```
1 if( pthread_mutex_trylock(&d->top_mutex) == 0)
2 {
3     l = d->queue[d->top % d->mem_size];
4     d->top++;
5     pthread_mutex_unlock (&d->top_mutex);
6 }
7 else{
8     l = abort_steal;
9 }
```

Listing 3.1: This excerpt taken from the code for the `de_steal` operation shows how the `pthread_mutex_trylock` function is used to avoid blocking when the mutex is already locked. If the function returns 0 the mutex is available otherwise the else clause is taken.

This allows the thread to handle such a situation accordingly; in the case of the steal operation an abort signal, and in the case of the pop operation an empty signal. An empty signal is produced here because the only situation where a `de_steal` operation will block a `de_pop_bottom` operation is when there is only one item remaining in the deque, which has already been claimed by the steal operation. The benefit of using this approach, rather than a `pthread_mutex_lock` based method, is that **dead-lock** is avoided.

The Circular Array

The deque makes use of a circular array which automatically grows and shrinks. This approach is memory efficient and intuitive. It also allows for the top counter to remain unchanged unless a steal operation occurs (in most cases), reducing the frequency of locks occurring. Figure 3.3 illustrates some of its mechanics. For instance figure 3.3v shows that the array is indeed circular, in that it “wraps around” when the final element is reached and there are elements un-used at the start.

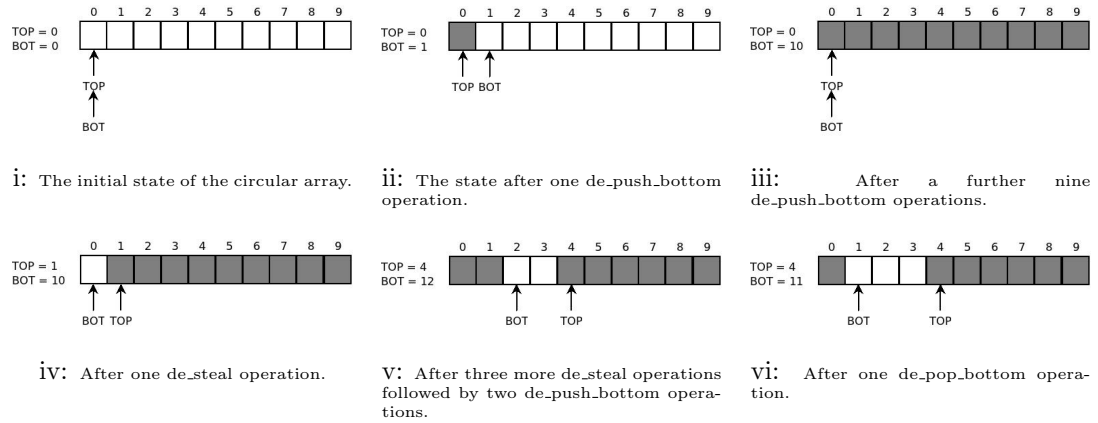


Figure 3.3: A demonstration of the a sequence of operations applied to the circular array data-structure. The ‘TOP’ field refers to the index of the top element of the deque and the ‘BOT’ field refers to the index of the bottom element plus one. White squares indicate an un-used element, shaded squares indicate elements which are in use.

The design relies on bottom and top counters (used to index the array), the memory-size (the number of memory slots allocated to the deque), and the size (the number of elements which are currently used). The size is computed by subtracting the top counter from the bottom counter. This is used to detect an empty array, and detect when the array needs to be re-sized. When the array is accessed, the correct element is indexed by computing $realindex = index \bmod memorysize$. For example, let the top counter equal fifteen and the memory-size equal ten, the $realindex$ of the top element is five. This allows for the top and bottom counters to exceed the memory size and still point to the correct index, without allowing access to memory outside the bounds of the array. It also maintains that the size can be computed correctly. When the array becomes empty the counters are re-set to zero.

Should the size exceed the memory-size, the array is automatically re-sized by double its current allocation. This is achieved by allocating a new array using `malloc`, copying the contents of the old array then freeing it, then assigning a pointer to the new array to the deque. Similarly, the array is shrunk by half when its size recedes to half of the memory-size using the same approach. The consequence of this design is that the `top_mutex` must be locked throughout any re-size operation, blocking any `de_steal` attempts until the mutex is un-locked. The benefits of this approach are that the deque is more scalable, and only memory that is likely to be used is allocated. Rather than some amount that is defined at compile-time.

3.3.2 The Work-Stealing Mechanism

The design described here utilises the deque implementation to produce an effective work-stealing scheme. The algorithm is so called ‘randomised’ because of the means to which a victim is selected.

- **ws_worker_thread:**

The function which is passed to `pthread_create`. This acts as the point of entry for each thread. It accepts a pointer to the deque which is associated with the given thread. This is passed in the form of a void pointer which has to be cast due to the interface provided by the pthreads library.

This function makes use of a do-while loop which breaks when there is no work

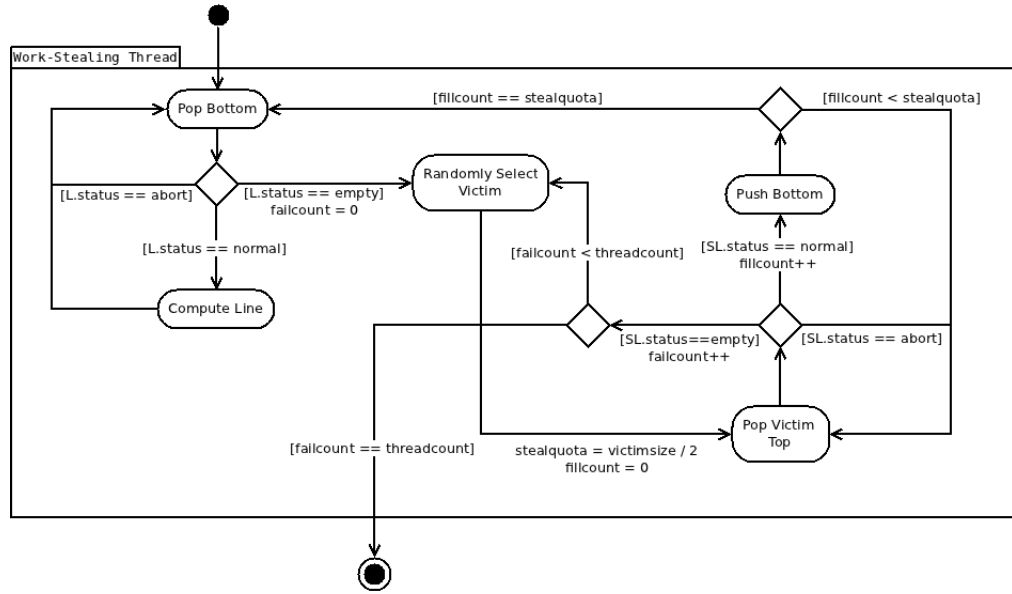


Figure 3.4: An activity diagram to describe the behaviour of a work-stealing thread which utilises the double ended queue outlined in section 3.3.1. L is the line popped from the threads deque. SL (stolen line) is the line popped from the victim's deque. The *victims* is the amount of work currently in the victim's deque. The *stealquota* is the amount of work items to be stolen. The *fillcount* is the number of items that have been stolen so far. The *failcount* is the frequency of failed steal attempts.

remaining in any worker-threads ready-deque. Each iteration of the loop sets the thread to work by calling the `ws_compute_deque` function, which returns when no work remains in the ready-deque. Then, the thread becomes a thief and must attempt to acquire work from other threads by calling `ws_become_thief`. If more work is acquired the thread becomes a worker and calls `ws_compute_deque` again. If no work is found the conditions of the loop are broken and the thread exits.

- **ws_compute_deque:**

Accepts a Deque pointer, i.e. to the Deque associated with the same thread. This function calls `de_pop.bottom` until it detects that the ready-deque is empty, at which point it returns the total number of work items completed for this call.

- **ws_become_thief:**

This operation randomly selects a thread to victimise. If it detects that there is no work available to steal it returns 0, otherwise it returns 1. It accepts a Deque pointer to the Deque associated with the same thread.

When a victim is selected the `ws_victimise` function is called.

In order to determine that no work is available, a set to mark each thread which is unsuccessfully victimised (referred to as the exclude set), as-well as a counter is maintained. The exclude set is initialised by adding the current thread. It is passed to `ws_random_deque` which uses the standard `c rand` function to generate a number between zero and the number of threads, used to access an array of deques indexed by thread id. The randomly selected deque is returned.

The function expects the exclude set to have at-least one element which is not a member.

- **ws_victimise:**

Accepts deque pointers for the thief and victim threads. This function is responsible for the attempted migration of half the victims workload, to the deque of the thief thread. It returns zero for an unsuccessful victimisation and one for success.

The initial line is acquired by calling `de_steal` on the victim. If the line has the status `LINE_EMPTY` zero is returned. Otherwise the victim has work to be stolen, so its work-count is evaluated and halved to determine how many work items the thief can attempt to steal. The initial line is pushed onto the thief's deque and the next line is stolen, providing the `fill_count` (the counter which tracks how much work has been stolen) has not exceeded the allotted amount of work to be stolen. Otherwise one is returned. The `fill_count` is then incremented. The line is checked for the `LINE_ABORT` status. If this is true `de_steal` is called again. Otherwise this entire process is repeated until the `LINE_EMPTY` status is received, or the allotted amount of work is stolen.

Problematic Characteristics

This scheme allows for multiple threads to perform work-stealing operations simultaneously. Due to the random nature of victim selection, some problematic situations may arise at run-time. These situations are likely to be detrimental to the overall balance of work load. They are described as follows:

- **Double Victimisation:**

Where two threads concurrently victimise the same thread. In this situation more than half of a victims deque is migrated. It can result, under certain circumstances, in almost all of the work in a victims deque being re-distributed in quick succession.

This is problematic because it can force unnecessary migration of work, which subsequently results in more frequent steal operations.

Figure 3.5i illustrates this phenomenon.

- **Victimisation of a Thief:**

When a thief is selecting a thread to victimise it does not discriminate between working threads and thief threads. Thus a 'sub-thief' can steal work from a thread which is already stealing work from a victim.

In this situation work is indirectly moved from a victim to the sub-thief, via another threads deque. The problem is the size of a thief thread's deque, does not equal the total number of stolen items until all items are migrated. As a result the sub-thief effectively steals less than half of the thief's eventual workload, as it evaluates its size before all items are transferred. This could result in an unfair re-distribution of work, and ultimately an unbalanced work load.

The diagram in figure 3.5ii represents such a situation.

Both problems could potentially be mitigated by making work-stealing operations atomic using mutex locks. In practice this involves implementing a mechanism where a thread can only take the state of either worker, thief, or victim at any given time, where thieves may only victimise workers. This would come at the cost of blocking execution of a thread in certain cases, and would add to the complexity of the algorithm.

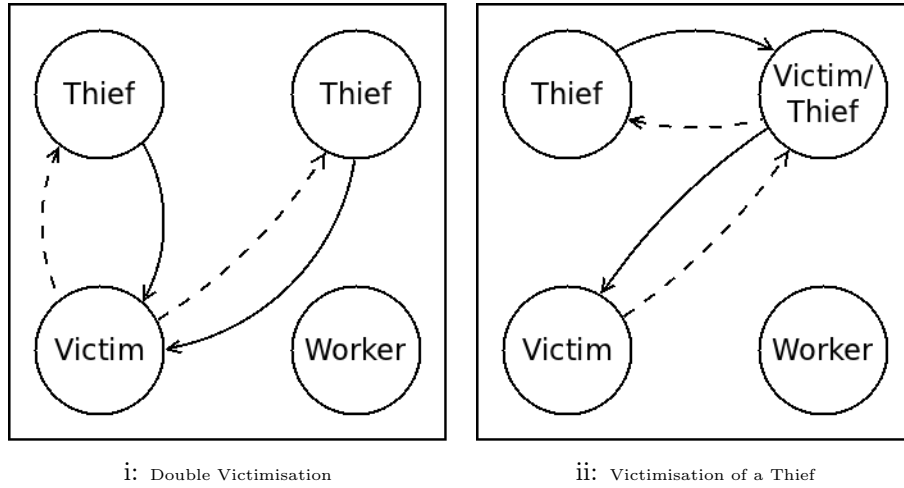


Figure 3.5: Visual representations of problematic situations. Dashed arrows show the flow of work items.

3.4 Additional Schemes

This section describes three additional modules which make use of different alternative approaches to that described in section 3.3. They interface with the Mandelbrot module in the same fashion as described in section 3.2

These schemes primarily serve as a benchmark for the randomized work stealing algorithm, but also serve to verify the Mandelbrot module.

A Sequential Algorithm

This is the simplest algorithm presented and utilizes no concurrency. It iterates each line (i.e. y index) of the raster plane in sequence, calling `compute_line` for each value of y . This module effectively implements a variation of the level set algorithm described in listing 2.1 by adding the outer for loop of the `compute_mandelbrot` function.

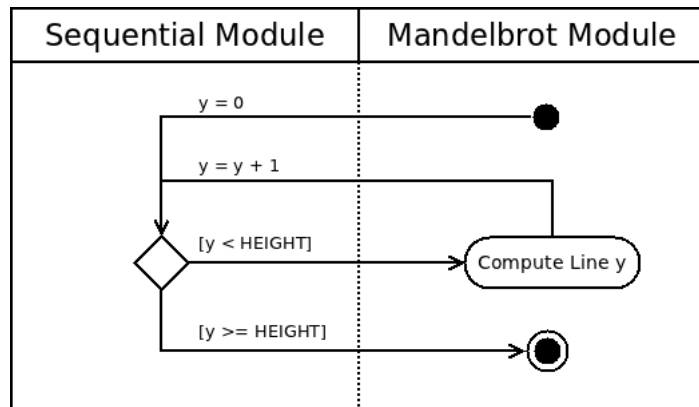


Figure 3.6: An activity diagram which illustrates the simplicity of the sequential scheme.

A Naïve Parallel Algorithm

This scheme is a simple parallel algorithm. It is so called naïve because of its comparable unbalanced work distribution to the a run-time work balanced algorithm. It

is an example of a typical scheme used for composing a parallel program, based on a sequential implementation.

The raster plane is divided into equal regions, which each thread computes independently. The regions are comprised of a set of contiguous y lines of the size $height/threadcount$. Each region is assigned as follows where $ystart$ is the first and $yend - 1$ is the last member of the region:

$$regionsize = height/threadcount$$

$$ystart = threadid * regionsize$$

$$yend = (threadid + 1) * regionsize$$

For the region with the thread id equal to $threadcount - 1$ (i.e. the thread which computes the lower region) $yend = height$. This is to ensure that all lines are assigned when $height$ is not integer divisible by $threadcount$. This functionality is omitted from figure 3.7 as not to muddy the simple, demonstrative purpose of the diagram.

Figure 3.7 illustrates how the design of the sequential version, documented in figure 3.6, is composed onto multiple threads to attain a parallel design. This is done by introducing a fork point, which starts the threads, and a join point which returns program control to the main execution flow.

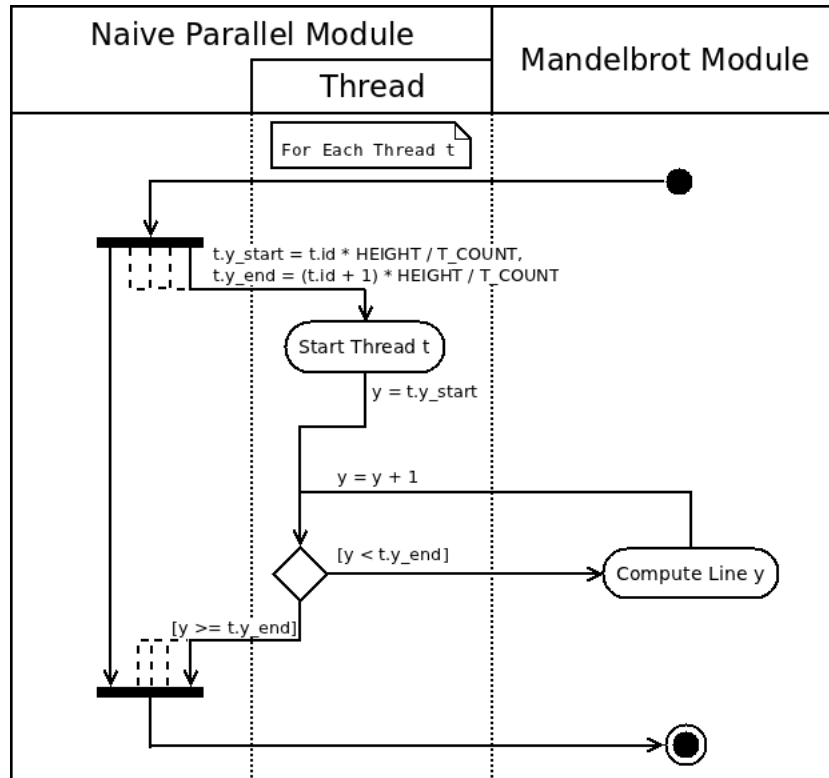


Figure 3.7: An activity diagram which illustrates the use of simple multi-threading employed by this scheme. It is worth noting that the dashed lines near the fork and join bars are added to illustrate that more than one thread executes simultaneously.

A Render-Thread Work-Stealing Algorithm

This scheme implements the approach described in the background research section 2.3.2. One monitor thread and three worker-threads are used.

The monitor thread is delegated the responsibility of controlling steal operations, and synchronising the thief and victim threads in order to complete this operation. Each render thread computes its work-load until it has no more to do. At this point it produces a thief signal which the monitor thread detects. A worker thread can be made a victim while it is working. It is allowed to finish its current line before becoming a victim.

The expected completion time of a render thread is evaluated to determine a victim. This means such a metric needs to be maintained, and regularly updated, in order to make this approach fair. The time taken to complete a work item is multiplied by the amount of work remaining for that thread. This is re-calculated after each line is computed.

Figure 3.8 shows a graphical representation of a typical, happy-day scenario steal operation. It highlights the amount of inter-thread communication which occurs in order to complete such an operation.

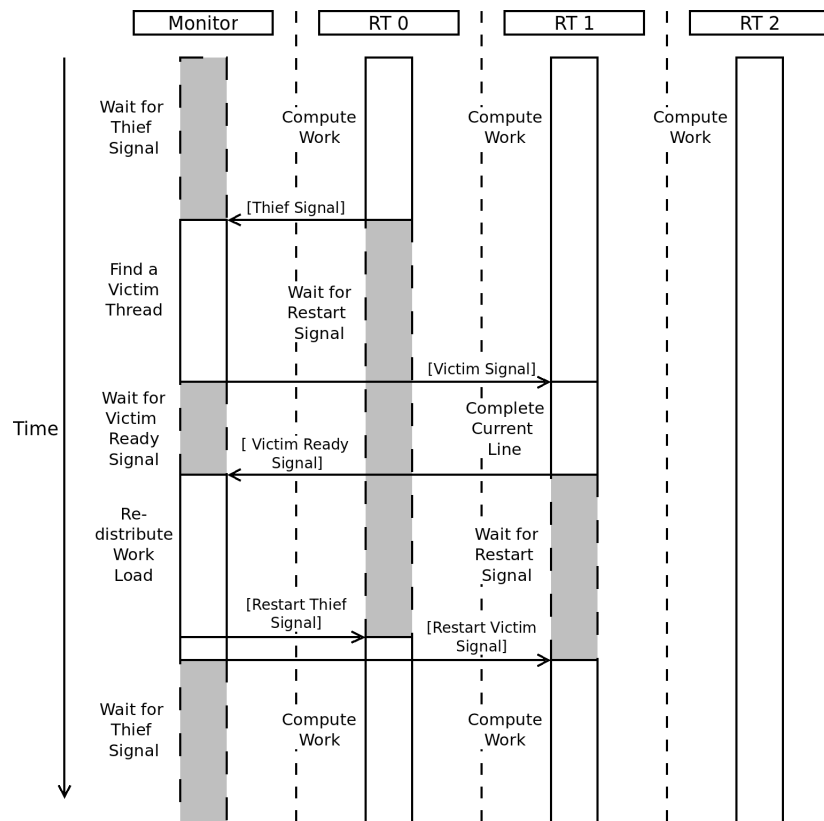


Figure 3.8: A diagram which illustrates a typical steal operation for the Render-Thread Algorithm. It illustrates how steal operations require thread synchronisation. Grey shaded boxes, with dashed borders, represent regions of time in which the corresponding thread must wait for a signal. White boxes indicate that a thread is performing some task. Arrows between threads indicates inter-thread communication.

This approach blocks working threads in order to achieve work-stealing. This comes with the added complication that such a scheme needs to be carefully designed in implementation.

Unfortunately, due to time constraints, there is no fully complete implementation for this scheme in place. It is worth noting that this area alone is a significant body of work, and represents a good portion of the future work this project could lead to.

3.5 Output

The system is capable of outputting PPM image files of the computed raster plane. This can be enabled using command line options described in section 3.7. Output is made optional because the process of producing the file is computationally time consuming, especially for large raster planes. It contributes to the code which cannot be parallelised.




Figure 3.9: Example images produced using the three implemented output modes. All three are produced using the Randomised Work-Stealing Algorithm detailed in section 3.3.

Output File Format

The specification for the PBM, PGM, and PPM formats is presented here [7].

Each file contains data for one image encoded in plain text. This consists of some header information, followed by an entry for each pixel.

The header is comprised of four fields. On line one the ‘magic number’ determines the colour style used. P1/P4 is used for monochrome (PBM), P2/P5 for grey-scale (PGM), and P3/P6 for RGB colour (PPM). On line two is the width and height as decimal numbers, separated by a space. On line three is the maximum value of a pixel. For the PBM format this line is omitted because it is redundant. Each line is separated with a carriage return.

Pixel data is separated by spaces and arranged in lines ending in a carriage return. Each line has exactly *width* entries and there are exactly *height* lines. For the PPM format each pixel has three entries representing red, green, and blue.

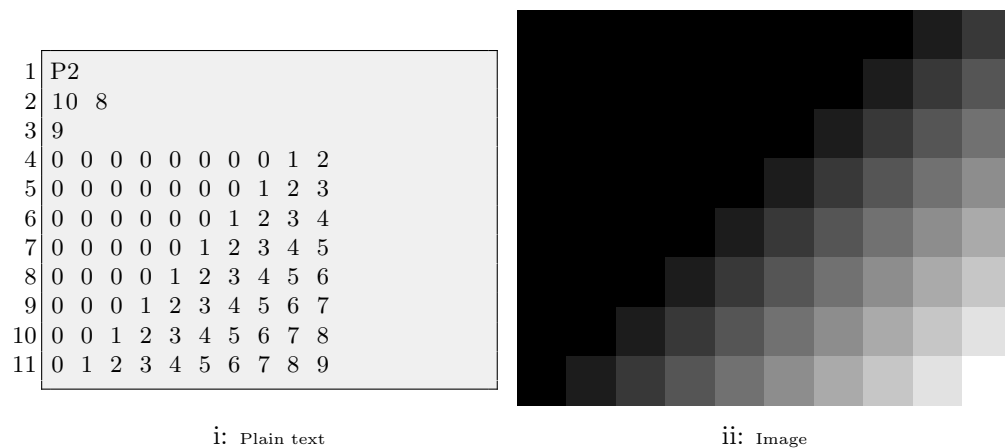


Figure 3.10: An example of the PGM file format and its corresponding image.

3.6 Run-Time Tracing

A simple tracing mechanism is implemented allowing for more detailed run-time analysis of the program. Any trace related code is not compiled by default, as it has a detrimental effect on performance. To enable this feature an option must be passed to ‘make’ when the project is built.

Related code is encased in preprocessor directives as follows:

```
1 /* Tracing for all modes */
2 #ifndef TRACE
3 ...
4 trace_event("mode 1: %",123);
5 #endif
6
7 /* Mode specific tracing code */
8 #if TRACE == 2
9 ...
10 trace_event("mode 2: %d",456);
11 #endif
```

The `trace_event` function prints a time-stamp in microseconds relative to the start of execution, followed by the desired message. A mutex is locked to ensure each output message is completed and written in sequence.

Currently two modes are in place. These can be built by passing the following arguments to ‘make’:

TRACE=1

A completion message is printed when the scheduling module has completed execution. This serves as a means of measuring the execution time of the program as a timestamp relative to the start in μ seconds.

```
1 499066: complete
```

TRACE=2

Traces thread activity in scheduler modules. All thread related traces specify a thread id using the `T_id` field. The following example shows a trace of the naïve parallel algorithm.

```
1 242: T_id 1 started
2 ...
3 754652: T_id 1 finished computing 1250 lines
4 ...
5 754804: complete
```

Tracing the Randomised Work-Stealing Algorithm

The trace format for the Randomised Work Stealing algorithm requires some explanation. It is detailed as follow:

```
1 1316: T_id 0 started
2 ...
3 396334: T_id 0 steal vi: 3, ws: 10
4 ...
5 1948938: T_id 3 ret-work fc: 1
6 ...
7 1960169: T_id 3 finished li: 3237, sc: 4, vc: 4
```

Listing 3.2: Examples of the Four Traced Events for the Randomised Work-Stealing Scheme

- **started:**

The thread specified by the T_id field has started. Line one of listing 3.2 shows an example.

- **steal:**

The thread specified by the T_id field has successfully performed a steal operation. The vi (victim) field indicates the thread id of the victim and ws (work stolen) field shows the quantity of work items migrated. Line three of listing 3.2 shows an example.

- **ret-work:**

The thread specified by the T_id field has finished stealing work and is returning to compute its stolen lines. The fc (fail count) field shows how many times the thread victimised another unsuccessfully. Line five of listing 3.2 shows an example.

- **finished:**

The thread specified by the T_id field has detected no available work and has finished. The li (lines) field shows the quantity of work items computed, the sc (steal count) field shows how many successful steal operations the thread carried out, and the vc (victimisation count) field indicates how many times the thread was victimised. Line seven of listing 3.2 shows an example.

Some examples of full run-time traces are available in the appendices section B.

3.7 User Interface

A basic command line interface is implemented which offers options to control the output of the program. When the program is given no arguments it computes the mandelbrot set with no output. The available options are described in the following paragraphs.

--outmode=<mode>

Determines the appearance of the output PPM file. The mechanism for producing output is described in section 3.5 and examples of each mode are shown in figure 3.9. There are three possible modes implemented:

- **greyscale** Gradient of black to white.
- **redscale** Gradient of black to red.
- **distribution** Each thread is represented by a different gradient.

This option must be used in order for output to occur.

--outfile=<file>

This option is used to specify the name of the output file. It must be used in conjunction with the 'outmode' option. If this option is omitted the output file takes the default name 'out.ppm'.

--help

Displays a usage message which describes the options detailed above.

Chapter 4

Discussion and Evaluation

4.1 Analysis of the Implemented Algorithms

This section discusses some observations made based on results gathered over a series of program runs. It also serves to assess the validity of the implemented schemes, ensuring that the code does what is expected.

The Test Setup

Here is a brief description of the environment used when gathering the results to support this section.

Test System

- CPU: Quad Core Intel Core I5 750 - 2.67GHz.
- Memory: 8GB DDR-3
- Operating System: Linux Mint 14 - Kernel 3.5.0-17 generic.
- C-Compiler: GCC Version 4.7.2 - Level 3 Optimisations

4.1.1 Performance Analysis

The following observations relate to performance of the implemented schemes. The collected performance data, for both figure 4.1 and 4.2, are calculated using an average time of five runs of the labelled version of the program per plot. These are presented in relation to the sequential version of the program.

The following equation is used to calculate the performance increases shown:

$$RelativePerformanceIncrease_a = \frac{AverageExecutionTime_a}{AverageExecutionTime_{sequential}}$$

- An average speed-up of 3.49 times is achieved for the randomised work-stealing algorithm. The naïve Parallel Algorithm achieves an average speed-up of 2.13 times.

This observation shows that by adding work-stealing techniques to a parallel scheme, which employs no load balancing, increases its performance by an average of 1.66 times.

Although both schemes make use of four threads, a speed up of four times is unlikely to be possible. This is due to factors such as thread scheduling overheads,

portions of code which cannot be parallelised, and imperfect load-balancing. This is associated with Amdahl's law, which states that such a speed-up is limited by the time taken for the sequential portion of the code to execute [28].

- According to figure 4.2, the performance of the Randomised work-stealing algorithm is unaffected by the maximum iterations parameter. The naïve Parallel Algorithm exhibits a performance decrease as the value is increased.

As the maximum iterations is increased the work required to compute a point in, or close to, the mandelbrot set is amplified. As a result the initially distributed work-load becomes further unbalanced.

These results suggest that the Randomised Work-Stealing algorithm exhibits good scalability, because of its load-balancing capabilities, when considering the maximum iterations setting. This point is also true for the plane size as shown in figure 4.1.

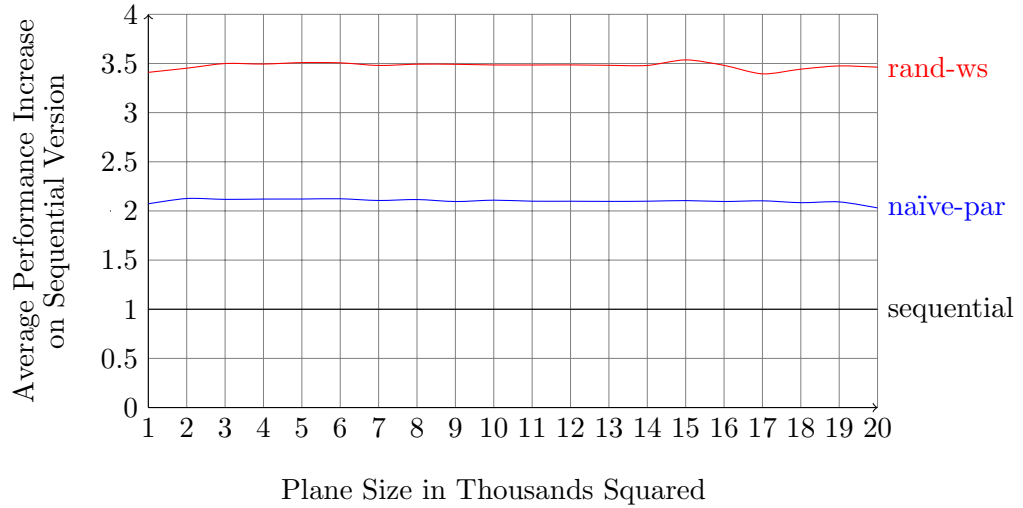


Figure 4.1: Average performance increase, relative to the sequential version of the program, against resolution of the raster-plane. The maximum iteration count is seventy for each sample.

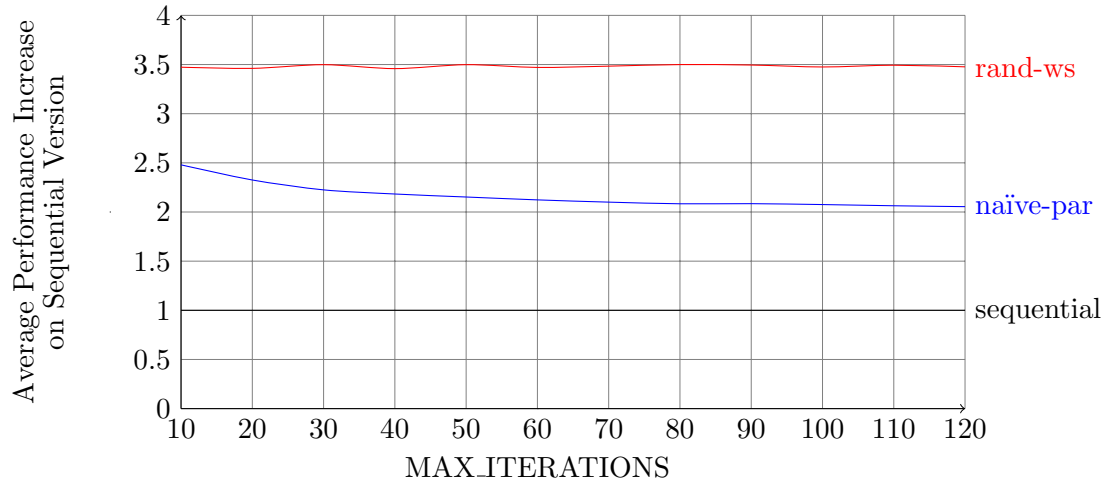


Figure 4.2: Average performance increase, relative to the sequential version of the program, against maximum number of iterations taken to determine the result of a single point. The raster-plane is ten-thousand squared in size for each sample.

4.1.2 Run-Time Analysis of the Randomised Work-Stealing Algorithm

The scatter plot shown in figure 4.3 and the bar charts shown in figure 4.4 bring to light some interesting properties of the algorithm. They also highlight some potential points of improvement.




Figure 4.3: A scatter plot where each point is a succesful steal operation. Each point is a mapping of time of occurance, and frequency of work items migrated in the operation. Points marked in black are collected from a single run of the program. Gray points are extracted from an additional four runs. All data is collected using trace mode two documented in section 3.6, where the raster plane is ten-thousand squared in dimension and the maximum iterations is set to seventy.

The following observations are based on the results presented in figure 4.3:

- As less work items are available, steal-operations become more frequent, in which fewer work items are migrated. This trend continues until no work items are available and completion is detected, as is suggested by the data.

The result is that a high frequency of steal operations, where very small amounts of work are migrated, occur as processing nears termination. The benefit of migrating such small work-loads could be seen to outweigh the cost of simply allowing the thread to continue computing the given region.

This could be alleviated by implementing a work migration threshold, which limits a steal operation to a minimum number of steal-able work. Should the projected work-load to be migrated fall under this threshold, the steal attempt would fail.

- A situation where multiple threads perform a steal operation in close time proximity arises often. Five clear examples of this are exposed at approximately 0.4,1.4,1.6,1.8, and 1.9 seconds. This outweighs the three instances in which a single steal operation occurs within any 0.1 second time-frame.

The following trace excerpt corresponds to the data shown in the plots marked in black. It details the three points plotted at approximately 1.6 seconds.


```
1 ...  
2 1625982: T_id 0 steal vi: 1, ws: 429  
3 1626009: T_id 3 steal vi: 1, ws: 214  
4 1628855: T_id 2 steal vi: 3, ws: 106  
5 ...
```

This example highlights some possible instances of the problematic situations detailed in section 3.3.

A double victimisation is present where both thread zero and three steal from thread one. Notably thread three migrates roughly half the work items thread zero does. This suggests that thread one's work load has been reduced by nearly three quarters in quick succession.

In addition, closely after, thread two victimises thread three. This suggests that victimisation of a thread has occurred, and again the transferred work-load is approximately half that stolen by the victim. It implies that work has indirectly been moved from thread zero, to thread two, via thread three.

This behaviour is a side effect of using a randomised approach to victimisation. Other approaches, for example evaluating work throughput metrics, could alleviate these problems.



i: Steal count and victimisation count per thread. ii: Percentage of total work-load computed per thread.

Figure 4.4: Sub-figure 4.4i shows a bar chart to display the average frequency of steal operations and the average frequency of victimisations, per thread id, over twenty test runs. Sub-figure 4.4ii shows a bar chart to show the average percentage of the raster-plane computed by each thread over the same twenty test runs.

The following observations are based on the bar charts presented in figure 4.4:

- Each thread is victimised on average 3.275 times over twenty program runs. The lowest average victimisation count is 2.9 (for thread zero) and the highest is 3.75 (for thread one). These results, which are visualised in figure 4.4i, suggest that randomly selecting a victim thread to steal from is a reasonably fair mechanism. It is worth noting that this is dependant on the qualities of the random number generator function. This implementation uses the ‘rand’ function provided the ‘stdlib’ library with a static seed. This is done to produce more repeatable behaviour. A more effective approach may be to produce a random function in which the seed is based on a dynamic system property, for instance the clock time.
- The threads which are initially assigned the upper and lowermost quarters of the raster-plane compute, on average, roughly three times more work items than the inner threads. This demonstrates that the algorithm implicitly balances work-load at run-time, and is shown in figure 4.4ii.

The behaviour exhibited here is expected, as the majority of the computational load lies in the central regions. Put more concisely; when a point is deemed a member of the set, exactly the maximum number of iterations will have occurred, thus involving more work. In general, the closer a point lies to the set, the more costly it becomes to process.

Thread ID	Completion Time (μ secs)
0	410922
1	3001814
2	2956266
3	432267

Table 4.1: The time taken from thread start to finish, for each thread in the naïve parallel algorithm to compute its assigned region. The data presented here is sourced from the trace presented in appendices figure B.3.

As the set has symmetry on the horizontal axis it is expected that each half

should represent an equal proportion of the work-load. This can be illustrated by analysing a trace of the naïve parallel algorithm. Table 4.1 shows that the central regions take longer to compute when no run-time load balancing is employed.

Split into four initial, equally sized regions the randomised work-stealing algorithm should ideally redistribute work to alleviate the exhibited behaviour of the naïve-parallel algorithm. For the outermost regions (i.e. threads zero and three) this is true, as can be seen in figure 4.4ii, where these ultimately compute approximately the same proportion of the plane. However, a discrepancy arises for the innermost threads (one and two) where thread one, on average, computes much fewer lines than its counterpart two.

This could be attributed to a number of aspects of the design. For instance, the phenomena of double victimisation and victimisation of a thief, compounded by an unsuitable pseudo random number generator, could be responsible for the unfairly balanced work load. Further investigation is required to determine a more evident cause.

4.2 Reflection

In general, the author feels that this project is a successful one. It has been an enlightening and thoroughly enjoyable experience to investigate the subject area. Here is an evaluation of each objective defined in the introduction section, as well as some other aspects of the project which are of notable merit.

Core Objectives

Background Research

Background research is presented in chapter 2. To start; a broad range of areas related to scheduling parallel algorithms is explored and explained. The work-stealing approach is introduced here, however the discussion is further developed in section 2.3. A concise explanation of the Mandelbrot set, and its surrounding themes, is offered between these. Finally, an overview of considered tools is added to show the extensive range of relevant development resources available.

The studied areas provide the reader with a sufficient overview of the themes explored in the subsequent chapters. To further bolster the readers understanding, a glossary of terms is presented in appendix A. This gives a brief definition for some of the major concepts discussed, which go slightly beyond the scope of this body of work.

These factors combined communicate a well researched background to the problem domain.

A Sequential Mandelbrot Set Algorithm

Although the sequential algorithm described in section 3.4 appears trivial, it relies almost entirely on the mandelbrot module described in section 3.2. This objective encompasses the implementation of this module, as this separation is derived from a sound design decision to encapsulate in order to reduce code complexity. It also improves code re-usability, of which is exploited extensively elsewhere throughout the project.

This body of work also demonstrates good application of the C programming language.

This objective is met with a high degree of quality, and forms a well-grounded foundation for the following objectives to build on.

A Naïve Parallel Mandelbrot Set Algorithm

A basic understanding of composing parallel programs is effectively demonstrated by completion of this objective. The described implementation in section 3.4 shows effective use of the pthreads library and its functionality.

A Random Work-Stealing Mandelbrot Set Algorithm

This objective forms a significant proportion of the work carried out during this project. It is the central focus of the project and exhibits more advanced programming and design techniques. This is described in great depth in section 3.3.

The implemented scheme demonstrates a strong grasp of concepts discovered through research of the subject area (see section 2.3), and utilises these in a proficient and competent fashion. The design is both elegant and simple, as-well as effective in practice. Some advanced techniques demonstrated here include; production of a non-blocking double ended queue, more advanced use of pthreads mutex locks, and design of a sophisticated work-load balancing algorithm.

Constructive criticism of the scheme is also offered, giving commentary on the implications of certain properties and possible remedies for problematic behaviour.

Overall, this portion of the project represents the majority of the technical and theoretical learning achievements made.

Analysis of the Implemented Algorithms

Technical analysis of the end product is offered in section 4.1. Observations made, based on data collected from traced run-time activity, are offered.

These provide some insight into the performance of the end product, as well as some explanation for undesirable behaviour.

This assesses how effective the randomised work-stealing algorithm is and identifies room for improvement. It also serves as validation to show that the product does what is intended, as well as to verify that it functions correctly.

This area proved challenging because of a lack of experience in such an exercise. This made making concise and well founded observations much more difficult. The work could be improved by producing better refined and well thought out test data sets.

Advanced Objectives

Graphical Output

The program produces output images in various styles. These are detailed in section 3.5. This proves that the product produces a correct approximation of the mandelbrot set (seeing is believing) and, in the case of distribution mode, gives evidence of the work-stealing algorithm's effectiveness.

The image format used is text based, thus inefficient and slow to write. The solution used to achieve this objective is basic, but serves the purpose.

A Render Thread Work-Stealing Mandelbrot Set Algorithm

This objective is incomplete as of the compilation of this report. Areas explored in an attempt to complete this additional objective include a design for the scheme, and a working but unstable implementation.

Work-Stealing Trace System

This objective is completed and provides an effective means for gathering and analysing run-time characteristics of an implemented scheme.

Although the objective is met a number of improvements could be made to allow tracing of a more detailed set of information.

Further Points of Reflection

Project Management

The major objectives for this project were completed comfortably in the time-scale provided. These achievements closely correspond with the planned time allocation described in the Gantt chart, which accompanies the detailed project proposal.

A major factor which streamlined the process is the decision to use a Git repository. It provided a vital means of tracking progress through regular commit messages, and the ability to easily revert back to older revisions. All the work related to this project can be found here: <https://github.com/mhawes/wstealmandel>. The available punch-card graph gives an idea of how frequently used the repository is, as well as an insight into the authors sleeping pattern over the past twelve weeks.

Professional Presentation of Material

All material is presented in a highly professional manner. This can be attributed to choosing the right tools for the job.

In particular, the decision to produce this entire report using L^AT_EX has been pivotal to the clarity and good structure exhibited. It has also provided a useful new skill set, which will be taken forward to future projects.

4.3 Further Work

Although the core project objectives have been completed, as well as the majority of the additional objectives, there are a number of interesting areas in which this work could be extended.

Some of these suggestions could form whole BSc Computer Science projects in their own right. The author hopes they may at least invoke some inspiration for those planning to embark on such a task.

Alternative Work-Stealing Schemes and Policies

To further investigate the work-stealing approach, configurations left un-explored by this report could be implemented. Areas which could be examined include:

- **Different Victimisation Policies:** Techniques for selecting a victim, for example choosing the thread which has the lowest work-throughput, could prove to more-fairly balance work-load. This would also serve as a point of comparison for the randomised approach.
- **Threshold Stealing:** By enforcing a minimum limit for the amount of work a thief can steal, a number of unnecessary steal operations could be avoided.
- **Alternative Deque Implementations:** The Deques implementation could be re-implemented using a different data structure design. For instance instead of a circular array, a linked list could be put in place.

- **Different Multi-Thread Designs:** Other approaches to parallel algorithm design could be explored to seek a more intuitive design. A fully working implementation of the Render-Thread scheme would be a good starting point for this line of enquiry.

Computing Julia Sets

The existing code in the Mandelbrot module could be fairly easily adapted to support raster-plane generation for Julia sets. The algorithm could accept a parameter to specify the value of z for the function described in equation 2.1. This would expose a huge set of new test cases for the randomised work-stealing algorithm implementation.

Investigate Specific Multi-Processor Architectures

Parallel computing platforms such as cluster computing, cloud computing, cellular architectures, amongst many others geared toward high performance computing, would be a relevant area for further study.

By focusing on a particular class of architecture the algorithm could be optimised, taking into consideration such properties as cache and memory layout, locality of related work-items, and processor locality.

An implementation which utilizes a cluster architecture, using message passing protocols, is of particular interest to the author.

Development of A Parallel Programming Library

The work-stealing techniques explored in this project could be applied to a vast number of programs. Producing a library which provides abstract programming directives, such as parallel data-structure iterators, or support for fully strict multi-threaded computations, could prove a useful general programming tool.

Expansion of the Capabilities of the Tracing Mechanism

A more concise tracing system could be implemented. The current mechanism provides only very limited run-time information of the implemented algorithms. Features to calculate statistics such as work item throughput per thread, detection of detrimental situations, and running tallies of thread status are desirable. These could provide vital debugging tools and could help to identify detrimental properties of the algorithm.

More Efficient Image Generation

Currently the output image takes a noticeably long time to write. This is due to the nature of the file format used.

The program could be adapted to use a dynamic graphics library, such as OpenGL, to display the image as it is generated. In addition, this could be used to illustrate the work-stealing scheme in action. As each thread completes a work item its corresponding line would appear on the screen in its assigned colour gradient. This would show how work items are migrated throughout runtime and serve to demonstrate the workings of the algorithm.

4.4 Conclusion

This study highlights a simple, yet effective, work-stealing scheme for computing a problem which exhibits unbalanced work-load when decomposed. It is applied to a parallel program which computes an approximation of the Mandelbrot set. This is compared and contrasted with both a sequential, and naïve parallel approach to the same problem. These deliverables are supported and guided by an extensive review of relevant literature.

The investigation finds that by applying the Randomised Work-Stealing Algorithm a significant speed-up is made over a similarly structured design, which neglects run-time load-balancing techniques. A further benefit is found in that the algorithm scales well with an increase in raster-place size and maximum iterations. Randomisation is also shown to be a fair method for victimisation. The analysis serves to identify some problematic situations in the form of double victimisation and victimisation of a thief, which could have an impact on the fairness of work-load redistribution.

All core objectives and all bar one advanced objectives are complete. The work done towards producing a Render-Thread Work-Stealing Algorithm, although incomplete in the time-scale, still offers some insight into an alternative approach to Work-Stealing.

The project as a whole is reflected upon in great depth and a good number of avenues for future work are laid out. Over-all the author considers the project a great success in the respect that the proposed work is complete, and a great deal of valuable knowledge has been acquired. The pursuit of this projects completion has allowed the author to hone his skills in many useful programming tools and techniques. The learning outcomes have not only been technical, good project and time management skills have also been gained.

The main and most valuable contribution is the implementation and design of a Randomised Work-Stealing algorithm. This demonstrates the technical ability of the author, as well as a strong understanding of programming and computer science theory. It also puts forward an intuitive and simple technique for improving a vast range of multi-threaded programs.

Chapter 5

Resources

References

- [1] About javaplot. ONLINE. <http://gnujavaplot.sourceforge.net/JavaPlot/About.html> accessed: 10th February 2013.
- [2] Fraqtive - mandelbrot family fractal generator. ONLINE. <http://fraqtive.mimec.org/> accessed: 5th February 2013.
- [3] The free opengl utility toolkit. ONLINE. <http://freeglut.sourceforge.net/> accessed: 4th March 2013.
- [4] Glut - the opengl utility toolkit. ONLINE. <http://www.opengl.org/resources/libraries/glut/> accessed: 4th March 2013.
- [5] Gnuplot homepage. ONLINE. <http://www.gnuplot.info/> accessed: 10th February 2013.
- [6] Gnuplot-iostream interface - c++. ONLINE. <http://www.stahlke.org/dan/gnuplot-iostream/> accessed: 10th February 2013.
- [7] Ppm format specification. ONLINE. <http://netpbm.sourceforge.net/doc/ppm.html> accessed: 10th February 2013.
- [8] S-net - home. ONLINE. <http://snet-home.org/> accessed: 17th April 2013.
- [9] Single assignment c – high productivity meets high-performance. ONLINE. <http://www.sac-home.org/> accessed: 17th April 2013.
- [10] Nimar S Arora, Robert D Blumofe, and Greg C Plaxton. Thread scheduling for multiprogrammed multiprocessors, 1998.
- [11] Blaise Barney. Posix threads programming. ONLINE, 2013. <https://computing.llnl.gov/tutorials/pthreads/> accessed: 10th February 2013.
- [12] Micheal F Barnsley, Robert L Devany, Benoit B Mandelbrot, Heinz-Otto Peitgen, and Dietmar Saupe Richard F Voss. *The Science of Fractal Images*. Springer-Verlag, 1988.
- [13] Robert D. Blumofe and Charles E. Leiserson. Scheduling multithreaded computations by work stealing. *Journal of the ACM*, 46(5):720–748, September 1999.
- [14] Robert D. Blumofe and Dionisios Papadopoulos. Hood: A user-level threads library for multiprogrammed multiprocessors, 1998.
- [15] David Chase and Yossi Lev. Dynamic circular work-stealing deque, 2005.
- [16] Peter Coad. *Java modeling in color with UML : enterprise components and process*. Prentice Hall PTR, Upper Saddle River, NJ, 1999.
- [17] Nathan Cohen. Fractals new era in military antenna design, 2005.

- [18] Nicolas Devillard. Gnuplot interfaces in ansi c. ONLINE. <http://ndevilla.free.fr/gnuplot/> accessed: 10th Febuary 2013.
- [19] David Feldman. *Chaos and Fractals: An Elementary Introduction*. OUP Oxford, 2012.
- [20] Vladimir Janjic. *Load Balancing of Irregular Parallel Applications on Heterogeneous Computing Environments*. PhD thesis, University of St. Andrews, 2012.
- [21] Bob Kuhn, Paul Petersen, and Eamonn OToole. Openmp versus threading in c/c++, 2000.
- [22] Jonathan Lifflander, Sriram Krishnamoorthy, and Laxmikant V Kale. Work stealing and persistence-based load balancers for iterative overdecomposed applications, 2012.
- [23] Benoit B Mandelbrot. How long is the coast of britain?, 1967.
- [24] Benoit B Mandelbrot. *Fractals: Form, Chance and Dimension*. W.H.Freeman and Co Ltd, 1977.
- [25] Benoit B Mandelbrot. *The Fractal Geometry of Nature*. W.H.Freeman and Co Ltd, 1983.
- [26] Jason McGuiness. Atomic code-generation techniques for micro-threaded risc architectures. Master’s thesis, University of Hertfordshire, 2006.
- [27] Steve R. Palmer and Mac Felsing. *A Practical Guide to Feature-Driven Development*. Pearson Education, 1st edition, 2001.
- [28] David A. Patterson and John L. Hennessy. *Computer Organization and Design*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 4th edition, 2008.
- [29] Michael M Resch, Alexander Schulz, Matthias S Muller, and Wolfgang E Nagel. *Tools for High Performance Computing 2009*. Springer-Verlag, 2010.
- [30] Scott Schneider, Christos D. Antonopoulos, and Dimitrios S. Nikolopoulos. Factory: An object-oriented parallel programming substrate for deep multiprocessors. *Proceedings of the 2005 International Conference on High Performance Computing and Communications*, pages 233–242, 2005.
- [31] Waide B Tristram. Investigating tools and techniques for improving software performance on multiprocessor computer systems. Master’s thesis, Rhodes University, 2011.

Chapter 6

Appendices

Appendix A

Glossary of Terms

Glossary

barrier A thread synchronisation technique in which a specified number of threads must reach an explicit point in execution before any may continue. 11

circular array An array in which the element indexed directly after the last element is effectively the first element, giving the effect of the data-structure “wrapping around”. 9

complex-number A number which is expressed in two parts; real and imaginary ($a + bi$). 5

condition-variable A thread synchronisation technique which allows conditional lock of a resource dependant on the state of some value. 11

critical-orbit The orbit of point 0 for a set of complex numbers. 7

dead-lock A situation in which two or more threads are waiting for each other to give up exclusive access to a resource required to continue execution. 18

exact self-similarity A property of a shape which is comprised of an exact copy of part of itself. 6

fractal A set which has a fractional dimension. 6

load-balancing A method in which work-load is distributed across multiple resources in an attempt to optimise utilisation of such resources. 5

locality The amount to which a value, or set of related values, is utilised by a resource. 5

mathematical-fractal Fractals which could exhibit the property of infinite scale invariance or quasi-self similarity. 6

monitor-thread A thread which performs administrative tasks such as maintaining worker-threads and controlling synchronisation. 9

multi-threaded algorithm A program which utilises more than one thread to achieve its desired result. 5

mutex In the context of pthreads, a mutex is a mechanism for acquiring exclusive lock on a resource. Short for mutual exclusion. 11

non-blocking A property of a multi-threaded algorithm in which shared resources are not locked, allowing threads to operate without being stopped. 9

quasi self-similarity A property of a shape which is comprised of an approximate copy of part of itself. 6

ready-deque A queue data-structure which allows access to both the top and bottom elements. 9

real-world-fractal Fractals which do not exhibit the property of scale invariance. 6

scale-invariance A property of a shape which is identical at magnifications of a common factor. 6

scheduling The process of determining the order in which parts of a program access system resources. 5

self-similarity A property of a shape which is comprised of part of itself. 6

steal-operation The process of re-distributing work-load from a victim to a thief. 9

thief A thread which has the work-load of some other thread re-assigned to its own. 8

thread A strand of execution which operates independently from the main flow of control in a process. 5

victim A thread which has a portion of its work-load re-assigned to some other thread. 8

work-sharing A processor which creates new work attempts to migrate it to another underutilised processor at creation time. 5

work-stealing A processor which is starved of work attempts to “steal” work from other processors. 5

worker-thread A thread which, in general, performs tasks that work towards achieving the purpose of the program. 9

Appendix B

Trace Output

The trace examples here are produced using the following constants in mandelbrot.h:

- HEIGHT = 10000
- WIDTH = 10000
- MAX_ITERATIONS = 70

```
1 2997274: complete
```

Listing B.1: An example trace produced by the ‘semandelbrot’ binary. Trace mode is one.

Other traces for mode one are omitted because they all take the same form.

```
1 4: Started
2 6311056: Finished
3 6311137: complete
```

Listing B.2: An example trace produced by the ‘semandelbrot’ binary. Trace mode is two.

```
1 264: T_id 2 started
2 270: T_id 1 started
3 292: T_id 0 started
4 315: T_id 3 started
5 411214: T_id 0 finished computing 2500 lines
6 432582: T_id 3 finished computing 2500 lines
7 2956530: T_id 2 finished computing 2500 lines
8 3002084: T_id 1 finished computing 2500 lines
9 3002311: complete
```

Listing B.3: An example trace produced by the ‘nsmandelbrot’ binary. Trace mode is two.

```
1 1289: T_id 1 started
2 1289: T_id 0 started
3 1318: T_id 2 started
4 1418: T_id 3 started
5 402418: T_id 3 steal vi: 2, ws: 818
6 402426: T_id 3 ret-work fc: 0
7 404182: T_id 0 steal vi: 2, ws: 408
8 404185: T_id 0 ret-work fc: 0
```

```

9 937474: T_id 2 steal vi: 3, ws: 277
10 937483: T_id 2 ret-work fc: 0
11 1090359: T_id 0 steal vi: 2, ws: 107
12 1090368: T_id 0 ret-work fc: 0
13 1362642: T_id 2 steal vi: 1, ws: 978
14 1362648: T_id 2 ret-work fc: 0
15 1375208: T_id 0 steal vi: 2, ws: 483
16 1375211: T_id 0 ret-work fc: 0
17 1544923: T_id 0 steal vi: 2, ws: 139
18 1544931: T_id 0 ret-work fc: 0
19 1625982: T_id 0 steal vi: 1, ws: 429
20 1625989: T_id 0 ret-work fc: 0
21 1626009: T_id 3 steal vi: 1, ws: 214
22 1626012: T_id 3 ret-work fc: 0
23 1628855: T_id 2 steal vi: 3, ws: 106
24 1628857: T_id 2 ret-work fc: 0
25 1805224: T_id 2 steal vi: 1, ws: 62
26 1805232: T_id 2 ret-work fc: 0
27 1815588: T_id 3 steal vi: 0, ws: 153
28 1815592: T_id 3 ret-work fc: 0
29 1922202: T_id 1 steal vi: 3, ws: 40
30 1922211: T_id 1 ret-work fc: 1
31 1922743: T_id 2 steal vi: 1, ws: 19
32 1922746: T_id 2 ret-work fc: 0
33 1946743: T_id 2 steal vi: 3, ws: 10
34 1946746: T_id 2 ret-work fc: 0
35 1948303: T_id 1 steal vi: 3, ws: 4
36 1948305: T_id 1 ret-work fc: 0
37 1954789: T_id 1 steal vi: 0, ws: 29
38 1954791: T_id 1 ret-work fc: 0
39 1955573: T_id 3 steal vi: 2, ws: 1
40 1955576: T_id 3 ret-work fc: 0
41 1958113: T_id 3 steal vi: 0, ws: 13
42 1958114: T_id 3 ret-work fc: 1
43 1958254: T_id 2 steal vi: 1, ws: 13
44 1958255: T_id 2 ret-work fc: 0
45 1978160: T_id 1 finished li: 855, sc: 3, vc: 6
46 1978439: T_id 0 finished li: 3872, sc: 5, vc: 3
47 1978493: T_id 2 finished li: 2010, sc: 7, vc: 6
48 1978682: T_id 3 finished li: 3262, sc: 5, vc: 5
49 1978762: complete

```

Listing B.4: An example trace produced by the ‘wsmandelbrot’ binary. Trace mode is two.

Appendix C

Source Code

The Mandelbrot Module

Listing C.1: mandelbrot.c

```
1 #include "mandelbrot.h"
2
3 static pixel_t plane[HEIGHT][WIDTH]; /* array to hold the generated image
   */
4
5 /* globals used to control the limits of the raster plane */
6 complex_t c_max; /* maximum value of c */
7 complex_t c_min; /* minimum value of c */
8 complex_t c_factor; /* value used to calculate space between each sample
   on the
9
10 raster plane */
11
12 /* ARGUMENT SETTINGS */
13 output_arg_t out_arg; /* the output mode
14 char outfile[21] = "out.ppm"; /* the output filename with a default of
   'out.ppm'
15
16 #ifdef TRACE
17 /* only use this for trace build */
18 unsigned long long start_time;
19 pthread_mutex_t trace_mutex;
20 #endif
21
22 /*
23
24 */
25
26 /* CODE STARTS HERE:
27
28 */
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1040
1041
1042
1043
1044
1045
1046
1047
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079
1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1210
1211
1212
1213
1214
1215
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2109
2110
2111
2112
2113
2114
2115
2116
2117
2118
2119
2120
2121
2122
2123
2124
2125
2126
2127
2128
2129
2130
2131
2132
2133
2134
2135
2136
2137
2138
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2209
2210
2211
2212
2213
2214
2215
2216
2217
2218
2219
2220
2221
2222
2223
2224
2225
2226
2227
2228
2229
2230
2231
2232
2233
2234
2235
2236
2237
2238
2239
2240
2241
2242
2243
2244
2245
2246
2247
2248
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2259
2260
2261
2262
2263
2264
2265
2266
2267
2268
2269
2270
2271
2272
2273
2274
2275
2276
2277
2278
2279
2280
2281
2282
2283
2284
2285
2286
2287
2288
2289
2290
2291
2292
2293
2294
2295
2296
2297
2298
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2309
2310
2311
2312
2313
2314
2315
2316
2317
2318
2319
2320
2321
2322
2323
2324
2325
2326
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2430
2431
2432
2433
2434
2435
2436
2437
2438
2439
2440
2441
2442
2443
2444
2445
2446
2447
2448
2449
2450
2451
2452
2453
2454
2455
2456
2457
2458
2459
2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480
2481
2482
2483
2484
2485
2486
2487
2488
2489
2490
2491
2492
2493
2494
2495
2496
2497
2498
2499
2500
2501
2502
2503
2504
2505
2506
2507
2508
2509
2510
2511
2512
2513
2514
2515
2516
2517
2518
2519
2520
2521
2522
2523
2524
2525
2526
2527
2528
2529
2530
2531
2532
2533
2534
2535
2536
2537
2538
2539
2540
2541
2542
2543
2544
2545
2546
2547
2548
2549
2550
2551
2552
2553
2554
2555
2556
2557
2558
2559
2560
2561
2562
2563
2564
2565
2566
2567
2568
2569
2570
2571
2572
2573
2574
2575
2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
2596
2597
2598
2599
2600
2601
2602
2603
2604
2605
2606
2607
2608
2609
2610
2611
2612
2613
2614
2615
2616
2617
2618
2619
2620
2621
2622
2623
2624
2625
2626
2627
2628
2629
2630
2631
2632
2633
2634

```

```

36     ws_initialise_threads();
37     ws_start_threads();
38
39
40 #if TRACE >= 1
41     trace_event( "complete\n");
42 #endif
43
44     perhaps_print();
45
46     pthread_exit(NULL);
47     return (0);
48 }
49
50 /*
51
52 */
51 void initialise( )
52 {
53     c_min.re = -2;
54     c_max.re = 2;
55
56     c_min.im = -2;
57     c_max.im = c_min.im + (c_max.re - c_min.re) * HEIGHT / WIDTH;
58
59     /* used to convert x, y of the raster plane into a complex number */
60     c_factor.re = (c_max.re - c_min.re) / (WIDTH - 1);
61     c_factor.im = (c_max.im - c_min.im) / (HEIGHT - 1);
62
63     #ifdef TRACE
64     pthread_mutex_init( &trace_mutex, NULL);
65     #endif
66 }
67
68 /*
69
70 */
69 inline double convert_y_coord( unsigned int y)
70 {
71     return c_max.im - y * c_factor.im;
72 }
73
74 /*
75
76 */
75 inline double convert_x_coord( unsigned int x)
76 {
77     return c_min.re + x * c_factor.re;
78 }
79
80 /*
81
82 */
81 /*
82  * returns the iteration count when not in the set. This produces a nice
83  * gradient effect.
84  */
85 char is_member(complex_t c)
86 {
87     char i;
88     complex_t z;
89
90     z.re = c.re;

```

```

91     z.im = c.im;
92
93     for(i = MAX_ITERATIONS; i >= 0; i--)
94     {
95         if( ((z.re * z.re) + (z.im * z.im)) > 4)
96         {
97             /* c is not a member of the set */
98             return i;
99         }
100
101         z = julia_func( z, c);
102     }
103
104     /* if it gets this far c is a member of the set */
105     return MAX_ITERATIONS;
106 }
107
108 /*

```

```

    */
109 /* returns true if outside the circle of radius 2 */
110 inline char is_outside_rad2( complex_t c)
111 {
112     /* if the number is outside the radius 2 we know it cant be in the
113     set */
114     /* note: no need for the sqrt function here comparing to 4 rather
115     than 2
116     * is faster */
117     return( ((c.re * c.re) + (c.im * c.im)) > 4 );
118 }
119 /*

```

```

    */
119 /* calculates the next value of: z = z^2 + c */
120 complex_t julia_func(complex_t z, complex_t c)
121 {
122     complex_t res;
123
124     res.im = 2 * z.re * z.im + c.im;
125     res.re = (z.re * z.re) - (z.im * z.im) + c.re;
126
127     return res;
128 }
129
130 /*

```

```

    */
131 /* Computes a line of the raster plane */
132 void compute_line( unsigned int y, char t_id)
133 {
134     complex_t c_cur;
135     unsigned int x;
136
137     c_cur.im = convert_y_coord( y);
138     for( x = 0; x < WIDTH; x++)
139     {
140         c_cur.re = convert_x_coord( x);
141
142         /* if the pixel is outside radius of two */
143         if( is_outside_rad2( c_cur)){
144             plane[y][x].t_id = WORKER_COUNT;
145             plane[y][x].val = PPM_BLACK; /* make the outer black */

```

```

146     }
147     else{ /* otherwise it could be in the mandelbrot set */
148         /* assign this thread id to the pixel */
149         plane[y][x].t_id = t_id;
150         /* assign the resulting value to the pixel */
151         plane[y][x].val= is_member( c_cur);
152     }
153
154     #ifdef TRACE
155     if( x == WIDTH / 2 && y > 5 && y < HEIGHT - 5 && plane[y][x].val
156         == 0){
157         trace_event("IDIOT line %d t_id %d\n",y,plane[y][x].t_id);
158     }
159     #endif
160 }
161
162 /*
163  */
164  /* UTIL FUNCTIONS:
165  */
166  /*
167  */
168  /* A simple function to parse command line arguments to allow some dynamic
169  * control from the user.
170  *
171  * NOTE: repeated arguments overwrite the previous and the rightmost
172  * repetition
173  * is used.
174  */
175 void handle_arguments( int argc, char *argv[])
176 {
177     int i;
178     char substr[11];
179
180     if( argc == 1){ /* in this case no arguments have been passed */
181         /* USE DEFAULT SETTINGS */
182         out_arg = OFF;
183     }
184     else{
185         for( i = 1; i < argc; i ++){
186             /* apply rules here */
187             /* OUTMODE RULES: */
188             if( strcmp(argv[i], "--outmode=greyscale") == 0){
189                 out_arg = GREYSCALE;
190             }
191             else if( strcmp(argv[i], "--outmode=redscale") == 0){
192                 out_arg = REDSCALE;
193             }
194             else if( strcmp(argv[i], "--outmode=distribution") == 0){
195                 out_arg = DISTRIBUTION;
196             }
197
198             /* OUTFILE RULE: */
199             else if( strncmp(argv[i], "--outfile=", 10) == 0){
200                 strncpy( outfile, argv[i] + 10, 15);
201             }
202
203             /* USAGE RULE: */

```

```

202         else if( strcmp(argv[i], "--help") == 0){
203             print_usage();
204             exit(0);
205         }
206
207         /* ERROR RULE: */
208         else{
209             print_usage();
210             exit(0);
211         }
212     }
213 }
214 }
215
216 /*
217
218 */
217 void print_usage()
218 {
219     printf( "Computing the Mandelbrot set in parallel using a
220             workstealing technique. \n"
221             "Author: Martin Hawes\n"
222             "\n"
223             "Usage:\n"
224             "    * No Arguments\n"
225             "        The raster-plane is computed but no output is done.\n"
226             "\n"
227             "    * --outmode=<mode>\n"
228             "        Enables output of the raster-plane in a number of ppm
229             configurations.\n"
230             "        Default filename is \"out.ppm\".\n"
231             "        modes:\n"
232             "            greyscale    - Gradient of Black to White.\n"
233             "            redscale     - Gradient of Black to Red.\n"
234             "            distribution - Gradient per thread. \n"
235             "\n"
236             "    * --outfile=<file >\n"
237             "        Specifies the name of the output file.\n"
238             "        Maximum of 20 characters.\n"
239             "\n"
240             "    * --help\n"
241             "        Displays this usage message.\n"
242             "\n\n"
243             "Constants:\n"
244             "    WORKER_COUNT:    %d\n"
245             "    MAX_ITERATIONS:  %d\n"
246             "    HEIGHT:          %d\n"
247             "    WIDTH:           %d\n"
248             "\n",
249             WORKER_COUNT, MAX_ITERATIONS, HEIGHT, WIDTH);
250 }
251
252 /*
253
254 */
251 /* Prints to output file depending on the arguments passed to the program
252 */
252 void perhaps_print()
253 {
254     switch( out_arg){
255         case GREYSCALE :
256             write_to_ppm_greyscale();
257             break;

```

```

258     case REDSCALE :
259         write_to_ppm_redscale();
260         break;
261     case DISTRIBUTION :
262         write_to_ppm_dist();
263     }
264 }
265
266 /*
267
268 */
269 void print_complex(complex_t *com)
270 {
271     printf("(%g + %g*i) ", com->re, com->im);
272 }
273
274 /*
275
276 */
277 void write_to_ppm_greyscale()
278 {
279     unsigned int x, y;
280
281     FILE *fp = fopen( outfile, "w+");
282     fprintf(fp, "P2\n%d %d\n%d\n", WIDTH, HEIGHT, MAX_ITERATIONS);
283
284     /* classic nested for loop approach */
285     for(y = 0; y < HEIGHT; ++y)
286     {
287         for(x = 0; x < WIDTH; ++x)
288         {
289             #ifdef TRACE
290             if( x == WIDTH / 2 && y > 5 && y < HEIGHT - 5 &&
291                plane[y][x].val == 0){
292                 trace_event("Missing Line %d t_id
293                             %d\n", y, plane[y][x].t_id);
294             }
295             #endif
296
297             fprintf(fp, "%i ", plane[y][x].val);
298         }
299         fprintf(fp, "\n");
300     }
301
302     fclose(fp);
303 }
304
305 /*
306
307 */
308 void write_to_ppm_redscale()
309 {
310     unsigned int x, y;
311
312     FILE *fp = fopen( outfile, "w+");
313     fprintf(fp, "P3\n%d %d\n%d\n", WIDTH, HEIGHT, MAX_ITERATIONS);
314
315     /* classic nested for loop approach */
316     for(y = 0; y < HEIGHT; ++y)
317     {
318         for(x = 0; x < WIDTH; ++x)
319         {
320             #ifdef TRACE

```

```

313         if( x == WIDTH / 2 && y > 5 && y < HEIGHT - 5 &&
314             plane[y][x].val == 0){
315             trace_event("Missing Line %d t_id
316                         %d\n",y,plane[y][x].t_id);
317         }
318         #endif
319         fprintf(fp, "%i 0 0 ", MAXITERATIONS - plane[y][x].val);
320     }
321     fprintf(fp, "\n");
322 }
323 fclose(fp);
324 }
325
326 /*
327
328 */
329 void write_to_ppm_dist()
330 {
331     unsigned int x, y;
332
333     FILE *fp = fopen( outfile, "w+");
334     fprintf(fp, "P3\n%d %d\n%d\n", WIDTH, HEIGHT, MAXITERATIONS);
335
336     /* classic nested for loop approach */
337     for(y = 0; y < HEIGHT; ++y)
338     {
339         for(x = 0; x < WIDTH; ++x)
340         {
341             #ifdef TRACE
342             if( x == WIDTH / 2 && y > 5 && y < HEIGHT - 5 &&
343                 plane[y][x].val == 0){
344                 trace_event("Missing Line %d t_id
345                             %d\n",y,plane[y][x].t_id);
346             }
347             #endif
348             fprintf(fp, "%i 0 %i ", get_red_val(plane[y][x]),
349                     get_blue_val(plane[y][x]));
350         }
351         fprintf(fp, "\n");
352     }
353     fclose(fp);
354 }
355
356 */
357 char get_red_val( pixel_t pix)
358 {
359     if( pix.t_id == WORKER.COUNT)
360     {
361         return pix.val;
362     }
363     if( pix.t_id % 4 == 0){
364         return pix.val;
365     }
366     else if( pix.t_id % 4 == 1){
367         return MAXITERATIONS - pix.val;

```

```

368
369     return 0;
370 }
371
372 /*
    */
373 char get_blue_val( pixel_t pix)
374 {
375     if( pix.t_id == WORKER.COUNT)
376     {
377         return pix.val;
378     }
379
380     if( pix.t_id % 4 == 2){
381         return pix.val;
382     }
383     else if( pix.t_id % 4 == 3){
384         return MAX_ITERATIONS - pix.val;
385     }
386
387     return 0;
388 }
389
390 /*
    */
391 /* Prints a timestamp followed by the message to the output */
392 void trace_event( const char *mess, ...)
393 {
394 #ifndef TRACE
395     va_list args;
396
397     unsigned long long time;
398     struct timeval tv_time;
399
400     gettimeofday(&tv_time, NULL);
401     time = (tv_time.tv_sec * 1000000 + tv_time.tv_usec) - start_time;
402
403     /* lock the trace mutex so that only one trace event can write at a
404        time */
405     pthread_mutex_lock(&trace_mutex);
406
407     printf("%llu: ", time);
408     va_start(args, mess);
409     vprintf(mess, args);
410     va_end(args);
411
412     pthread_mutex_unlock(&trace_mutex);
413 #endif
414 }
415 /*
    */
416 inline void trace_start()
417 {
418 #ifndef TRACE
419     struct timeval tv_start;
420     gettimeofday(&tv_start, NULL);
421     start_time = tv_start.tv_sec * 1000000 + tv_start.tv_usec;
422 #endif
423 }

```

Listing C.2: mandelbrot.h

```

1 #ifndef MANDELBROT_H
2 #define MANDELBROT_H
3
4 #include <stdarg.h>
5 #include <math.h>
6 #include <time.h>
7 #include <string.h>
8 #include <stdio.h>
9 #include <stdlib.h>
10
11 /* depending on the make rule applied this attaches a different scheduler
   */
12 #ifdef rtmandel
13 #include "rtworksteal.h"
14 #endif
15 #ifdef wsmandel
16 #include "worksteal.h"
17 #endif
18 #ifdef semandel
19 #include "sequential.h"
20 #endif
21 #ifdef nsmandel
22 #include "noscheduling.h"
23 #endif
24 #ifdef inmandel
25 #include "interleaved.h"
26 #endif
27
28
29 /* dimensions of the raster plane */
30 #define HEIGHT 10000
31 #define WIDTH 10000
32
33 #define PPMBLACK 0 /* the value of black in a grey-scale ppm file */
34
35 #define MAXITERATIONS 70 /* between: 1-127 */
36
37 /* struct for nicely containing Complex numbers */
38 typedef struct complex_t{
39     double re, im;
40 } complex_t;
41
42 typedef struct pixel_t{
43     char t_id, val;
44 } pixel_t;
45
46 /* ARGUMENT ENUMS */
47 typedef enum {OFF, GREYSCALE, REDSCALE, DISTRIBUTION} output_arg_t;
48
49 /*
50
51 */
52 /* function declarations */
53
54 void initialise ( );
55 inline double convert_x_coord ( unsigned int);
56 inline double convert_y_coord ( unsigned int);
57 inline char is_outside_rad2 ( complex_t);
58 char is_member ( complex_t);
59 complex_t julia_func ( complex_t, complex_t);
60 void computeline ( unsigned int, char);

```

```
60 /* UTIL FUNCTIONS */
61 void handle_arguments      ( int , char *[] );
62 void print_usage           ();
63
64 void perhaps_print         ();
65
66 void print_complex         ( complex_t * );
67
68 void write_to_ppm_dist     ();
69 void write_to_ppm_greyscale ();
70 void write_to_ppm_redscale ();
71
72 char get_red_val           ( pixel_t );
73 char get_blue_val         ( pixel_t );
74
75 inline float fast_inv_sqrt ( float );
76
77 void trace_event           ( const char * , ... );
78 inline void trace_start    ();
79 #endif /* MANDELBROT.H */
```

The Randomised Work-Stealing Scheduling Module

Listing C.3: worksteal.c

```
1 #include "worksteal.h"
2
3 static deque_t dequees[WORKER_COUNT]; /* Set of dequeues to fill. One per
4 thead. */
5
6 pthread_t threads[WORKER_COUNT]; /* set of threads to execute the dequeues
7 */
8
9 #ifdef TRACE
10 int steal_count[WORKER_COUNT];
11 int victim_count[WORKER_COUNT];
12 #endif
13
14 /*
15 */
16 /* thread function */
17 /*
18 */
19
20 void *ws_worker_thread( void *t_deq)
21 {
22     deque_t *deq = (deque_t *) t_deq;
23     char stealable = 1;
24     int work_count = 0;
25
26     #if TRACE >= 2
27         trace_event("T_id %d started\n", deq->t_id);
28     #endif
29
30     do
31     {
32         /* This is where the work is done. */
33         work_count += ws_compute_deque( deq);
34
35         /* After this the thread turns into a thief */
36         stealable = ws_become_thief( deq);
37     } while(stealable == 1);
38
39     #if TRACE >= 2
40         trace_event("T_id %d finished li: %d, sc: %d, vc: %d\n", deq->t_id,
41             work_count, steal_count[deq->t_id], victim_count[deq->t_id]);
42     #endif
43
44     pthread_exit(NULL);
45 }
46
47 /*
48 */
49 /* this initialises the dequeues which are passed to the threads. */
50 void ws_initialise_threads()
51 {
52     char i;
53
54     /* initialise all dequeues */
55     for (i = 0; i < WORKER_COUNT; i++)
56     {
57         de_initialise( &dequees[i], i);
58     }
59 }
```

```

51
52     ws_distribute_lines();
53
54     /* set a seed for the rand function */
55     //      srand( 938672);
56
57     /* seed from the current usec val */
58     //      struct timeval tv_seed;
59     //      gettimeofday(&tv_seed, NULL);
60     //      srand( tv_seed.tv_usec);
61 }
62
63 /*

```

```

64 */
65 void ws_distribute_lines()
66 {
67     unsigned int y, distribution = HEIGHT / WORKER_COUNT;
68     char i = -1;
69     line_t line;
70
71     line.status = LINE_NORMAL;
72
73     /* initialise dequeues with work before starting */
74     for( y = 0; y < HEIGHT ; y++)
75     {
76         line.y = y;
77
78         if( y % distribution == 0){
79             i++;
80         }
81
82         /* distribute the current line y to one of the dequeues */
83         de_push_bottom( &dequeues[i], line);
84     }
85 }
86 /*

```

```

87 */
88 void ws_start_threads()
89 {
90     pthread_attr_t attr;
91     void *status;
92     char i;
93
94     /* Initialize and set thread detached attribute */
95     pthread_attr_init(&attr);
96     pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
97
98     /* start threads */
99     for(i=0; i<WORKER_COUNT; i++) {
100         pthread_create(&threads[i], &attr, ws_worker_thread, (void
101             *)&dequeues[i]);
102     }
103
104     /* join point */
105     for(i=0; i<WORKER_COUNT; i++) {
106         pthread_join(threads[i], &status);
107     }
108
109     /* Free attribute and wait for the other threads */
110     pthread_attr_destroy( &attr);

```

```

109 }
110
111 /*
    */
112 unsigned int ws_compute_deque( deque_t *deq)
113 {
114     // printf( "T %d BECAME WORKER\n", deq->t_id );
115     unsigned int work_count = 0;
116     line_t line_cur;
117
118     while(1)
119     {
120         line_cur = de_pop_bottom( deq );
121
122         if( line_cur.status == LINE_EMPTY ){
123             break;
124         }
125
126         compute_line( line_cur.y, deq->t_id );
127
128         work_count++;
129     }
130
131     return work_count;
132 }
133
134 /*
    */
135 /* returns 1 if found work and needs to return to worker mode.
136    * returns 0 if no work is found in the entire network and the thread can
137       finish.
    */
138 char ws_become_thief( deque_t *deq)
139 {
140     int i;
141     char result = 0, ex_count = 1;
142     deque_t *victim;
143     char exclude_set[WORKER_COUNT];
144
145     /* initialise the exclude set */
146     for( i = 0; i < WORKER_COUNT; i++ )
147     {
148         exclude_set[i] = 0;
149     }
150     /* add this threads deque to the exclude set */
151     exclude_set[deq->t_id] = 1;
152
153     /* while no work has been stolen.
154        * in the case that no work is available 0 is
155        * returned and the loop is broken.
156        */
157     while( result == 0 )
158     {
159         /* if the exclude set is not exhausted get another victim at
160            random */
161         if( ex_count < WORKER_COUNT ){
162             victim = ws_random_deque( exclude_set );
163             result = ws_victimise( deq, victim );
164         }
165         else{
166             return 0;

```

```

166     }
167
168     /* add the unsuccessful victim to the exclude set */
169     if( result == 0){
170         exclude_set[victim->t_id] = 1;
171         ex_count++;
172     }
173 }
174
175 #if TRACE >= 2
176     /* increment the victimisation count */
177     victim_count[victim->t_id]++;
178     /* report the failure count as the thread returns to do the work it
179        stole */
179     trace_event("T_id %d ret-work fc: %d\n", deq->t_id, ex_count - 1);
180 #endif
181     return 1;
182 }
183
184 /*
185
186     */
185 /* returns 1 if work has been stolen and placed in the deque ready for
186    computing
187    * returns 0 if no work is found at this victim.
188    */
188 char ws_victimise( deque_t *deq, deque_t *victim)
189 {
190     char result = 0;
191     line_t line;
192     int steal_size;
193     int fill_count = 0;
194
195     /* evaluate victim size to work out how much to steal */
196     steal_size = (victim->bot - victim->top) / 2;
197
198     /* if there is nothing worth stealing return 0 */
199     if( steal_size == 0){
200         return 0;
201     }
202
203     /* steal a line */
204     line = de_steal( victim);
205
206     /* check if the victim is empty first */
207     if( line.status == LINE.EMPTY ){
208         return 0;
209     }
210
211     /* loop until we get an empty line or the right amount of work is
212        stolen */
212     while( line.status != LINE.EMPTY)
213     {
214         /* if we have a normal line push it onto this threads deque */
215         if( line.status == LINE.NORMAL){
216             de_push_bottom( deq, line);
217             /* when we have stolen the right amount of work from a victim
218                give up */
218             if( fill_count < steal_size){
219                 line = de_steal( victim);
220                 fill_count++;
221             }
222             else{

```

```

223 #if TRACE >= 2
224         steal_count[deq->t_id]++;
225         trace_event("T_id %d steal vi: %d, ws: %d\n", deq->t_id,
                    victim->t_id, fill_count);
226 #endif
227         /* push the stray line */
228         return 1;
229     }
230 }
231
232     /* if the thread was blocked try again */
233     if( line.status == LINEABORT){
234         line = de_steal( victim);
235     }
236 }
237
238     return 0;
239 }
240
241 /*
242
243     */
244 /*
245     * Generates a random number between 0 and WORKER_COUNT (not inclusive)
246     * and returns a deque that is NOT the same as any in the exclude_set.
247     * NOTE: this function takes no responsibility for an exclude set that is
248     *       fully
249     *       set.
250     */
251 deque_t *ws_random_deque( char exclude_set[WORKER_COUNT])
252 {
253     int i, j;
254
255     do
256     {
257         i = rand() % WORKER_COUNT;
258     } while(exclude_set[i] == 1);
259
260     return &deques[i];
261 }

```

Listing C.4: worksteal.h

```

1 #ifndef WORKSTEALH
2 #define WORKSTEALH
3
4 #define WORKER_COUNT 4          /* the total number of worker threads */
5
6 #include <pthread.h>
7 #include "deque.h"
8
9 /* HAS TO BE INCLUDED LAST!!! */
10 #include "mandelbrot.h"
11
12 void *ws_worker_thread          ( void* );
13
14 char ws_become_thief            ( deque_t * );
15 char ws_victimise               ( deque_t *, deque_t * );
16
17 void ws_distribute_lines        ( );
18 deque_t *ws_random_deque        ( char[WORKER_COUNT] );
19 unsigned int ws_compute_deque( deque_t *deq );
20
21 /*
22  */
23 void ws_initialise_threads      ( );
24 void ws_start_threads           ( );
25
26 /*
27  */
28 #endif /* WORKSTEALH */

```

The Deque

Listing C.5: deque.c

```
1 #include "deque.h"
2
3 /*
4  */
5 void de_initialise( deque_t *d, char thread_id)
6 {
7     d->t_id = thread_id;
8     d->mem_size = INIT_MEM_SIZE;
9
10    d->top = 0;    /* points to the NEXT position of top */
11    d->bot = 0;    /* points to the bototm of the queue */
12
13    /* allocates the initial block of memory */
14    d->queue = (line_t *) malloc(d->mem_size * sizeof(line_t));
15
16    /* initialise of EMPTY and ABORT vals. */
17    empty.status = LINE_EMPTY;
18    abort_steal.status = LINE_ABORT;
19
20    /* mutex for top access */
21    pthread_mutex_init( &d->top_mutex, NULL);
22 }
23
24 /*
25  */
26 /* Pushes a line onto the bottom of the queue.
27  */
28 void de_push_bottom( deque_t *d, line_t line)
29 {
30     int size = d->bot - d->top;
31
32     de_attempt_grow( d, size);
33
34     line.status = LINE_NORMAL;
35     d->queue[ d->bot % d->mem_size ] = line;
36
37     d->bot++;
38 }
39
40 /*
41  */
42 /* Takes the bottom member of the queue and increments the bottom counter.
43  * This function has a chance to shrink the size of the queue.
44  */
45 line_t de_pop_bottom( deque_t *d)
46 {
47     line_t l;
48     d->bot--;
49
50     int size = d->bot - d->top;
51     if( size < 0){
```

```

52     d->bot = d->top;
53     return empty;
54 }
55
56 l = d->queue[d->bot % d->mem.size];
57
58 if( size > 0){
59     /* in this case we want to attempt to shrink the array */
60     de_attempt_shrink( d, size);
61
62     return l;
63 }
64
65 /* in this case bottom is also top and lock needs to be checked for */
66 if( pthread_mutex_trylock(&d->top_mutex) == 0){
67     pthread_mutex_unlock (&d->top_mutex);
68     return l;
69 }
70 else {
71     l = empty;
72     d->bot = d->top + 1;
73 }
74
75 return l;
76 }
77
78 /*

```

```

    */
79 /* Takes the top member of the queue.
80 */
81 line_t de_steal( deque_t *d)
82 {
83     int size = d->bot - d->top;
84     line_t l;
85
86     /* in this case we only have the bottom member therefore do not
87     * want to steal.
88     */
89     if(size <= 0){
90         return empty;
91     }
92
93     /* mutex for top element needs to be enforced here. */
94     /* If the mutex is locked an abort signal line is returned */
95     if( pthread_mutex_trylock(&d->top_mutex) == 0)
96     {
97         l = d->queue[d->top % d->mem.size];
98         d->top++;
99         pthread_mutex_unlock (&d->top_mutex);
100     }
101     else{
102         l = abort_steal;
103     }
104
105     return l;
106 }
107
108 /*

```

```

    */
109 /*

```

```

110  */
111  char de_attempt_shrink( deque_t *d, int size)
112  {
113      int mem_s = d->mem_size;
114      if( size <= mem_s / 2 && mem_s > INIT_MEM_SIZE){
115          d->mem_size = mem_s / 2;
116          de_re_allocate( d, mem_s);
117          return 1;
118      }
119      return 0;
120  }
121  /*
122  */
123  char de_attempt_grow( deque_t *d, int size)
124  {
125      int mem_s = d->mem_size;
126      if(size >= mem_s){
127          d->mem_size = mem_s * 2;
128          de_re_allocate( d, mem_s);
129      }
130  }
131  /*
132  */
133  /* This function grows the array to the current value of mem_size.
134   * It does so by allocating a new array and copying the results from
135   * the old array over.
136   */
137  void de_re_allocate( deque_t *d, int old_size)
138  {
139      pthread_mutex_lock(&d->top_mutex);
140
141      int i, j = 0, size = d->bot - d->top;
142      line_t *temp_q = (line_t *)malloc(d->mem_size * sizeof(line_t));
143
144      for( i = d->top; i < d->bot; i++)
145      {
146          temp_q[i % d->mem_size] = d->queue[i % old_size];
147      }
148
149      de_free_queue( d);
150      d->queue = temp_q;
151
152      pthread_mutex_unlock (&d->top_mutex);
153  }
154  /*
155  */
156  */
157  void de_free_queue( deque_t *d)
158  {
159      free(d->queue);
160  }
161  /*
162  */
163  /* UTIL FUNCTIONS: */

```

```

163 /*
    */
164 void de_print_deque( deque_t *d)
165 {
166     int size = d->bot - d->top;
167     int i;
168     line_t l;
169
170     printf("Deque for thread id: %d\n", d->t_id);
171     printf("  Bot: %d Top: %d\n", d->bot, d->top);
172     printf("  Size:          %d\n", size);
173     printf("  Size in Memory: %d\n", d->mem_size);
174
175     printf("  Members:\n");
176     if( size >= 0){
177         for(i = d->mem_size - 1; i >= 0 ; i--)
178         {
179             l = d->queue[i];
180             if( i < 10){ printf( "      i:%d ",i);}
181             else{ printf( "      i:%d",i);}
182             printf( " [y:%d]", l.y);
183             if(i == d->bot % d->mem_size){ printf(" <- bot");}
184             if(i == d->top % d->mem_size){ printf(" <- top");}
185             printf("\n");
186         }
187     }
188     else{
189         printf("    Deque empty\n");
190     }
191     printf("\n");
192 }
193
194 void de_test_deque( deque_t *d)
195 {
196     line_t l;
197
198     /* fill up the deque */
199     int i;
200     for( i = 0; i < 700; i++)
201     {
202         l.y = i;
203
204         if(i % 2 == 0){ de_pop_bottom( d); }
205         if(i % 3 == 0){ de_steal( d); }
206
207         de_push_bottom( d, l);
208     }
209
210
211
212     while(1)
213     {
214         de_print_deque( d);
215
216         l = de_pop_bottom( d);
217         if( l.status == LINE_EMPTY){
218             printf("GOT EMPTY SIGNAL\n\n");
219             break;
220         }
221     }
222 }

```

Listing C.6: deque.h

```

1  /*
2   * Implementation of a double ended queue of line structs which is
3   * non-blocking.
4   * The deque is circular. (see paper)....
5   *
6   * The deque grows and shrinks accordingly.
7   */
8  #ifndef DEQUE_H
9  #define DEQUE_H
10
11 #include <stdlib.h>
12 #include <stdio.h>
13 #include <pthread.h>
14
15 #define LINE_NORMAL 0
16 #define LINE_EMPTY 1
17 #define LINE_ABORT 2
18
19 #define INIT_MEM_SIZE 10 /* the initial number of allocated Line slots */
20
21 typedef struct line_t{
22     char status;
23     unsigned int y;
24 } line_t;
25
26 typedef struct deque_t{
27     char t_id;
28
29     int mem_size;
30     int top, bot;
31
32     line_t *queue;
33
34     pthread_mutex_t top_mutex;
35 } deque_t;
36
37 static line_t empty, abort_steal;
38
39 /* function defs */
40
41 void de_initialise      ( deque_t *, char);
42
43 /* ----- */
44 line_t de_steal        ( deque_t *);
45 void de_push_bottom    ( deque_t *, line_t);
46 line_t de_pop_bottom   ( deque_t *);
47 /* ----- */
48
49 char de_attempt_shrink ( deque_t *, int);
50 char de_attempt_grow   ( deque_t *, int);
51 void de_re_allocate    ( deque_t *, int);
52 void de_free_queue     ( deque_t *d);
53
54 /* UTIL FUNCTIONS */
55 void de_print_deque( deque_t *);
56
57 #endif /* DEQUE_H */

```

The Sequential Module

Listing C.7: sequential.c

```
1 #include "sequential.h"
2
3 void ws_initialise_threads() {}
4
5 void ws_start_threads()
6 {
7     unsigned int y;
8
9     #if TRACE >= 2
10         trace_event("Started\n");
11     #endif
12
13     for( y = 0; y < HEIGHT; y++)
14     {
15         compute_line( y, 0);
16     }
17
18     #if TRACE >= 2
19         trace_event("Finished\n");
20     #endif
21 }
```

Listing C.8: sequential.h

```
1 #ifndef SEQUENTIAL_H
2 #define SEQUENTIAL_H
3
4 #define WORKER_COUNT 1
5
6 /* HAS TO BE INCLUDED LAST!!! */
7 #include "mandelbrot.h"
8
9 void ws_initialise_threads    ();
10 void ws_start_threads        ();
11
12 #endif /* SEQUENTIAL_H */
```

The Naïve Parallel Module

Listing C.9: noscheduling.c

```
1 #include "noscheduling.h"
2
3 pthread_t threads[WORKER_COUNT]; /* set of threads to execute */
4 thread_info_t infos[WORKER_COUNT];
5
6 /*
7  *
8  */
9 void *na_worker_thread( void* tia)
10 {
11     int y;
12     thread_info_t *ti = (thread_info_t *) tia;
13     na_initialise( ti);
14     #if TRACE >= 2
15         trace_event("T_id %d started\n", ti->t_id);
16     #endif
17     for(y = ti->start_y; y < ti->end_y; y++)
18     {
19         compute_line( y, ti->t_id);
20     }
21
22     #if TRACE >= 2
23         trace_event("T_id %d finished computing %d lines\n", ti->t_id,
24                     ti->end_y - ti->start_y);
25     #endif
26     pthread_exit(NULL);
27 }
28
29 /*
30  *
31  */
32 /* calculates the values of start_y and end_y.
33  */
34 void na_initialise( thread_info_t *ti)
35 {
36     char id = ti->t_id;
37     unsigned int distribution = HEIGHT / WORKER_COUNT;
38
39     ti->start_y = id * distribution;
40
41     if( id != WORKER_COUNT - 1){
42         ti->end_y = (id + 1) * distribution;
43     }
44     else{
45         ti->end_y = HEIGHT;
46     }
47 }
48
49 /*
50  *
51  */
52 void ws_initialise_threads()
```

```

53     thread_info_t ti;
54
55     /* give the threads ids */
56     for( i = 0; i < WORKER_COUNT; i++)
57     {
58         ti.t_id = i;
59         infos[i] = ti;
60     }
61 }
62
63 /*
64
65 */
64 void ws_start_threads()
65 {
66     pthread_attr_t attr;
67     void *status;
68     char i;
69
70     /* Initialize and set thread detached attribute */
71     pthread_attr_init(&attr);
72     pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
73
74     /* start threads */
75     for(i=0; i<WORKER_COUNT; i++) {
76         pthread_create(&threads[i], &attr, na_worker_thread, (void
77             *)&infos[i]);
78     }
79
80     /* join point */
81     for(i=0; i<WORKER_COUNT; i++) {
82         pthread_join(threads[i], &status);
83     }
84
85     /* Free attribute */
86     pthread_attr_destroy( &attr);
87 }

```

Listing C.10: noscheduling.h

```
1 #ifndef NOSCHEDULING_H
2 #define NOSCHEDULING_H
3
4 #define WORKER_COUNT 4      /* the total number of worker threads */
5
6
7 #include <pthread.h>
8
9 /* HAS TO BE INCLUDED LAST!!! */
10 #include "mandelbrot.h"
11
12 typedef struct thread_info_t{
13     char t_id;
14     unsigned int start_y, end_y;
15 } thread_info_t;
16
17 void *na_worker_thread      ( void*);
18 void na_initialise          ( thread_info_t*);
19
20 /*
21
22 */
21
22 void ws_initialise_threads  ();
23 void ws_start_threads      ();
24
25 /*
26
27 */
26
27 #endif /* NOSCHEDULING_H */
```

The Render-Thread Module

Listing C.11: rtworksteal.c

```
1 #include "rtworksteal.h"
2
3 pthread_t threads[WORKER_COUNT]; /* set of threads to execute the dequeues
   */
4 pthread_t monitor;
5 pthread_attr_t attr;
6 thread_info_t thread_infos[WORKER_COUNT];
7
8 pthread_mutex_t thief_mut; /* used to make sure we have 1 thief at a
   time */
9
10 /* used to stop/restart the victim/thief threads FIXME*/
11 /*pthread_mutex_t vi_stop_mut;
12 pthread_mutex_t th_stop_mut;
13 pthread_cond_t vi_stop_cond;
14 pthread_cond_t th_stop_cond;
15 */
16
17 pthread_mutex_t mon_wait_mut;
18 pthread_cond_t mon_wait_cond;
19 pthread_mutex_t thief_sig_mut;
20 pthread_cond_t thief_sig_cond;
21
22 pthread_mutex_t the_one_thief_mut;
23
24 pthread_barrier_t stop_bar;
25
26 char finish_sig;
27
28 /*
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1040
1041
1042
1043
1044
1045
1046
1047
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079
1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1210
1211
1212
1213
1214
1215
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2109
2110
2111
2112
2113
2114
2115
2116
2117
2118
2119
2120
2121
2122
2123
2124
2125
2126
2127
2128
2129
2130
2131
2132
2133
2134
2135
2136
2137
2138
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2209
2210
2211
2212
2213
2214
2215
2216
2217
2218
2219
2220
2221
2222
2223
2224
2225
2226
2227
2228
2229
2230
2231
2232
2233
2234
2235
2236
2237
2238
2239
2240
2241
2242
2243
2244
2245
2246
2247
2248
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2259
2260
2261
2262
2263
2264
2265
2266
2267
2268
2269
2270
2271
2272
2273
2274
2275
2276
2277
2278
2279
2280
2281
2282
2283
2284
2285
2286
2287
2288
2289
2290
2291
2292
2293
2294
2295
2296
2297
2298
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2309
2310
2311
2312
2313
2314
2315
2316
2317
2318
2319
2320
2321
2322
2323
2324
2325
2326
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2430
2431
2432
2433
2434
2435
2436
2437
2438
2439
2440
2441
2442
2443
2444
2445
2446
2447
2448
2449
2450
2451
2452
2453
2454
2455
2456
2457
2458
2459
2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480
2481
2482
2483
2484
2485
2486
2487
2488
2489
2490
2491
2492
2493
2494
2495
2496
2497
2498
2499
2500
2501
2502
2503
2504
2505
2506
2507
2508
2509
2510
2511
2512
2513
2514
2515
2516
2517
2518
2519
2520
2521
2522
2523
2524
2525
2526
2527
2528
2529
2530
2531
2532
2533
2534
2535
2536
2537
2538
2539
2540
2541
2542
2543
2544
2545
2546
2547
2548
2549
2550
2551
2552
2553
2554
2555
2556
2557
2558
2559
2560
2561
2562
2563
2564
2565
2566
2567
2568
2569
2570
2571
2572
2573
2574
2575
2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
2596
2597
2598
2599
2600
2601
2602
2603
2604
2605
2606
2607
2608
2609
2610
2611
2612
2613
2614
2615
2616
2617
2618
2619
2620
2621
2622
2623
2624
2625
2626
2627
2628
2629
2630
2631
2632
2633
2634
2635
2636
2637
2638
2639
2640
2641
2642
2643

```

```

53
54     pthread_exit(NULL);
55 }
56
57 /*
58
59 */
58 void *rt_monitor_thread( void *null)
59 {
60     printf("Monitor started\n");
61
62     void *status;
63
64     char i, is_work = 1;
65     char alive_count = WORKER_COUNT;    // used to catch straglers
66
67     thread_info_t *thief;
68     thread_info_t *victim;
69
70
71     /* start threads */
72     for(i = 0; i < WORKER_COUNT; i++) {
73         pthread_create(&threads[i], &attr, rt_render_thread, (void
74             *)&thread_infos[i]);
75     }
76
77     do
78     {
79         printf("Monitor waiting for thief\n");
80
81         thief = rt_wait_for_thief();
82
83         if( thief == NULL){
84             rt_broadcast_finished();
85             printf("Monitor detected no work");
86             break;
87         }
88
89         printf("Monitor found thief: %d\n", thief->t_id);
90
91         victim = rt_find_victim();
92
93         if( victim == NULL){
94             rt_broadcast_finished();
95             printf("Monitor detected no work after no victim\n");
96             break;
97         }
98
99         printf("Monitor found victim: %d\n", victim->t_id);
100
101         /* Wait for the finish_line mutex to be unlocked */
102         pthread_mutex_lock( &mon_wait_mut);
103         printf("Monitor waiting for victim to finish line\n");
104         victim->sta_v_sig = 1;
105         rt_print_workload( victim);
106         while( victim->sta_v_sig)
107         {
108             pthread_cond_wait(&mon_wait_cond, &mon_wait_mut);
109         }
110         pthread_mutex_unlock( &mon_wait_mut);
111         printf("Monitor got signal from victim\n");
112     }

```

```

113     rt_print_status(thief);
114     rt_print_workload( thief);
115     rt_print_status(victim);
116     rt_print_workload( victim);
117 */
118
119     pthread_mutex_lock( &thief->stop_mut);
120     pthread_mutex_lock( &victim->stop_mut);
121
122     rt_distribute( thief , victim);
123
124     /* restart the threads */
125     pthread_cond_signal( &victim->stop_cond);
126     pthread_cond_signal( &thief->stop_cond);
127     pthread_mutex_unlock( &victim->stop_mut);
128     pthread_mutex_unlock( &thief->stop_mut);
129
130     printf("Monitor finished distribution and restarted threads\n");
131
132     /* let threads continue */
133     //pthread_barrier_wait(&stop_bar);
134
135 } while( is_work == 1);
136
137 printf("Monitor is cleaning up the stragglers\n");
138 /* wait for stranglers before joining */
139 do
140 {
141     //thief = rt_wait_for_thief();
142     pthread_cond_signal( &thief_sig_cond);
143
144     for( i = 0; i < WORKER_COUNT; i++)
145     {
146         pthread_cond_signal( &thread_infos[i].stop_cond);
147         pthread_cond_signal( &thread_infos[i].stop_cond);
148
149         if( thread_infos[i].sta_finished){
150             printf("%d\n", alive_count);
151             alive_count--;
152         }
153     }
154 } while( alive_count > 0);
155
156
157 /* join point */
158 for(i = 0; i < WORKER_COUNT; i++) {
159     pthread_join(threads[i], &status);
160 }
161
162 printf("Monitor finished\n");
163
164 pthread_exit(NULL);
165 }
166
167 /*
168
169 */
170 void rt_distribute( thread_info_t *thief , thread_info_t *victim)
171 {
172     unsigned int block;
173     unsigned int victim_count;
174
175     /* calculate the work count of the victim */

```

```

174     victim_count = victim->end - victim->curr;
175
176     /*
177     printf("VICTIM b: ");
178     rt_print_workload( victim);
179     printf("THIEF b: ");
180     rt_print_workload( thief);
181     */
182
183     /* Steal exactly half of the victims work */
184     if(victim->curr < victim->end){
185         block = victim_count / 2;
186
187         /* re-assign the work */
188         thief->end = victim->end;
189         thief->curr = victim->curr + block + 1;
190         victim->end = victim->curr + block;
191     }
192
193     // thief->status = THREAD_WORKING;
194     // victim->status = THREAD_WORKING;
195
196     /*
197     printf( "steal-count: %u\n", block);
198     printf("VICTIM: ");
199     rt_print_workload( victim);
200     printf("THIEF: ");
201     rt_print_workload( thief);
202     */
203 }
204
205 /*
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232

```

```

    /*
    Sets all threads to finished status */
void rt_broadcast_finished()
{
    char i;

    //printf("FINISH BROADCAST\n");

    for( i = 0; i < WORKER_COUNT; i++)
    {
        //rt_print_status(&thread_infos[i]);
        //rt_print_workload(&thread_infos[i]);

        finish_sig = 1;
        //pthread_barrier_wait(&thread_infos[i].line_bar);

        thread_infos[i].status = WORK_FINISHED;

        /*
        pthread_mutex_unlock( &th_stop_mut);
        pthread_mutex_unlock( &vi_stop_mut);
        pthread_mutex_unlock( &thief_mut);
        */
    }

    // pthread_barrier_wait(&stop_bar);

```



```

233 }
234
235 /*
    */
236 /* Cycles through all the thread infos until it finds one that has a
    complete
237 * signal.
238 *
239 * Returns the thread_info of the complete thread.
240 */
241 thread_info_t *rt_wait_for_thief()
242 {
243     char i;
244     thread_info_t *ti = NULL;
245
246     do
247     {
248         for( i = 0; i < WORKER_COUNT; i++)
249         {
250             if( thread_infos[i].status == THIEF_SIG){
251                 thread_infos[i].status = IS_THIEF;
252                 ti = &thread_infos[i];
253                 pthread_cond_signal( &thief_sig_cond);
254             }
255         }
256     } while( ti == NULL && finish_sig != 1); /* repeat until no
        work-detected */
257
258     return ti;
259 }
260
261 /*
    */
262 /* Finds the thread with the highest estimated complete time */
263 thread_info_t *rt_find_victim()
264 {
265     char i, count = 0;
266
267     thread_info_t *result = &thread_infos[0]; /* take t_id 0 as the
        default */
268
269     for(i = 0; i < WORKER_COUNT; i++)
270     {
271         if( thread_infos[i].curr >= thread_infos[i].end){
272             count++;
273         }
274         else if( thread_infos[i].estimated_complete >
            result->estimated_complete &&
            thread_infos[i].sta_working)
275         {
276             thread_infos[i].sta_working = 0;
277             result = &thread_infos[i];
278         }
279     }
280
281
282     /* if I have counted WORKER_COUNT worth of threads that are empty it
        means
283     * we can complete
284     */
285     if( count == WORKER_COUNT){
286         return NULL;

```

```

287     }
288
289     /* set the vitim sig and let the victim finish its current line */
290     result->status = VICTIM_SIG;
291
292     return result;
293 }
294
295 /*

```

```

    */
296 unsigned int rt_compute_work( thread_info_t *ti)
297 {
298     unsigned int i;
299     unsigned int work_count = 0;
300
301     struct timeval tv_start, tv_end;
302
303     printf("t_id %d started doing work\n", ti->t_id);
304     ti->sta_working = 0;
305
306     while( ti->curr <= ti->end && finish_sig != 1 && ti->sta_v_sig == 0)
307     {
308         i = ti->curr;
309
310         gettimeofday(&tv_start, NULL);
311
312         compute_line( i, ti->t_id);
313         ti->curr++;
314         work_count++;
315
316         gettimeofday(&tv_end, NULL);
317
318         /* THIS IS NOT NICE! FIXME is there a better way to do this? */
319         rt_update_estimate( ti, ((tv_end.tv_sec -
320                                     tv_start.tv_sec)*1000000 +
321                                     tv_end.tv_usec - tv_start.tv_usec) / 100);
322     }
323
324     /* When this thread gets the victim signal it needs to stop working
325     now */
326     if( ti->sta_v_sig)
327     {
328         ti->sta_v_sig = 0;
329         rt_become_victim( ti);
330     }
331
332     return work_count;
333 }
334
335 /*

```

```

    */
334 void rt_update_estimate( thread_info_t *ti, unsigned long time)
335 {
336     unsigned int work_count = ti->end - ti->curr;
337
338     ti->estimated_complete = work_count * time;
339 }
340
341 /*

```

```

    */

```

```

342 /* this initialises the threads. */
343 void ws_initialise_threads()
344 {
345     char i;
346     unsigned int block = HEIGHT / WORKER_COUNT;
347     unsigned int prev_end = 0;
348
349     /* initialise the global mutexes */
350     /* pthread_mutex_init( &thief_mut, NULL);
351     pthread_mutex_init( &th_stop_mut, NULL);
352     pthread_mutex_init( &vi_stop_mut, NULL);
353     pthread_cond_init( &th_stop_cond, NULL);
354     pthread_cond_init( &vi_stop_cond, NULL);*/
355
356     pthread_mutex_init( &mon_wait_mut, NULL);
357     pthread_cond_init( &mon_wait_cond, NULL);
358     pthread_mutex_init( &thief_sig_mut, NULL);
359     pthread_cond_init( &thief_sig_cond, NULL);
360
361     pthread_mutex_init( &the_one_thief_mut, NULL);
362
363     finish_sig = 0;
364
365     // pthread_barrier_init(&stop_bar, NULL, 3);
366
367     for( i = 0; i < WORKER_COUNT; i++)
368     {
369         thread_infos[i].t_id = i;
370         thread_infos[i].estimated_complete = MAX_ESTIMATE;
371         thread_infos[i].status = THREAD_WORKING;
372
373         thread_infos[i].sta_t_sig = 0;
374         thread_infos[i].sta_v_sig = 0;
375         thread_infos[i].sta_stop = 0;
376         thread_infos[i].sta_working = 0;
377
378         pthread_mutex_init( &thread_infos[i].stop_mut, NULL);
379         pthread_cond_init( &thread_infos[i].stop_cond, NULL);
380
381         // pthread_barrier_init( &thread_infos[i].line_bar, NULL, 2);
382
383         /* if this is the last thread distribute right upto the last line
384         This is in-case the height isn't divisible by the worker count
385         */
386         thread_infos[i].curr = prev_end;
387         if( i == WORKER_COUNT - 1){
388             thread_infos[i].end = HEIGHT - 1;
389         }
390         else{
391             prev_end = block * (i + 1);
392             thread_infos[i].end = prev_end;
393             prev_end++;
394         }
395     }
396
397     /*
398
399     */
400 void ws_start_threads()
401 {
402     void *status;

```

```

402  /* Initialize and set thread detached attribute */
403  pthread_attr_init(&attr);
404  pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
405
406  /* start threads */
407  /* NB: the monitor thread has responsibility for starting the workers
    */
408  pthread_create(&monitor, &attr, rt_monitor_thread, (void *)NULL);
409
410  /* join the monitor thread */
411  pthread_join(monitor, &status);
412
413  /* Free attribute and wait for the other threads */
414  pthread_attr_destroy(&attr);
415 }
416
417 /*


---


    */
418 /* Issues the work complete signal and waits until its status is changed
    from
419  * THIEF_SIG.
420 */
421 void rt_become_thief( thread_info_t *ti)
422 {
423     pthread_mutex_lock( &the_one_thief_mut);
424     printf("t_id %d locked the_one_thief_mut\n", ti->t_id);
425
426     /* Attempt to lock the thief_sig mutex.
427      * Wait until chosen as a thief by the monitor.
428     */
429
430     pthread_mutex_lock( &thief_sig_mut);
431     printf("t_id %d locked thief_sig_mut\n", ti->t_id);
432
433     ti->status = THIEF_SIG;
434     ti->sta_t_sig = 1;
435
436     /* Set the estimated time to 0.
437      * This stops the monitor from picking the thief as a victim
438     */
439     ti->estimated_complete = 0;
440
441     while (ti->sta_t_sig)
442     {
443         pthread_cond_wait(&thief_sig_cond, &thief_sig_mut);
444         ti->sta_t_sig = 0;
445     }
446     pthread_mutex_unlock( &thief_sig_mut);
447     printf("t_id %d un-locked thief_sig_mut\n", ti->t_id);
448
449     /* make the thread wait while the distribution is being done */
450     pthread_mutex_lock( &ti->stop_mut);
451     printf("t_id %d stopped (T)\n", ti->t_id);
452     ti->status = THIEF_SIG;
453     ti->sta_stop = 1;
454     while (ti->sta_stop)
455     {
456         pthread_cond_wait(&ti->stop_cond, &ti->stop_mut);
457         ti->sta_stop = 0;
458     }
459
460     /* set the status to working */

```

```

461     ti->status = THREAD_WORKING;
462
463     pthread_mutex_unlock( &ti->stop-mut);
464     pthread_mutex_unlock( &the_one_thief-mut);
465     printf("t_id %d un-locked the_one_thief-mut\n", ti->t_id);
466     printf("t_id %d re-started (T)\n", ti->t_id);
467
468
469
470     /* Stop mutex makes sure we are stopped here. */
471     /* pthread_mutex_lock( &th-stop-mut);
472     while (ti->status == IS_THIEF)
473     {
474         pthread_cond_wait(&th-stop-cond, &th-stop-mut);
475     }
476
477     pthread_mutex_unlock( &th-stop-mut); */
478
479     /* wait for thief, victim, and monitor to call this function */
480 //    pthread_barrier_wait(&stop-bar);
481
482     /* Allow waiting thieves to become a thief */
483 //    pthread_mutex_unlock( &thief-mut);
484 }
485
486 /*

```

```

    */
487 /*
488 */
489 void rt_become-victim( thread-info_t *ti)
490 {
491     /* make the thread wait while the distribution is being done */
492     pthread_mutex_lock( &ti->stop-mut);
493
494     /* signal the monitor that it can start distributing */
495     pthread_cond_signal(&mon-wait-cond);
496
497     printf("t_id %d stopped (V)\n", ti->t_id);
498     ti->status = IS_VICTIM;
499     ti->sta-stop = 1;
500     while (ti->sta-stop)
501     {
502         pthread_cond_wait(&ti->stop-cond, &ti->stop-mut);
503         ti->sta-stop = 0;
504     }
505     /* set the status to working */
506     ti->status = THREAD_WORKING;
507
508     pthread_mutex_unlock( &ti->stop-mut);
509     printf("t_id %d re-started (V)\n", ti->t_id);
510
511 //    pthread_cond_signal( &ti->finish-line-cond);
512
513     /* when monitor and victim reach this barrier the line has ended */
514 //    pthread_barrier_wait(&ti->line-bar);
515
516     /* make this wait for the stop mutex to be unlocked */
517     /* pthread_mutex_lock( &vi-stop-mut);
518
519     while (ti->status == IS_VICTIM)
520     {
521         pthread_cond_wait(&vi-stop-cond, &vi-stop-mut);

```

```

522     }
523
524     pthread_mutex_unlock( &vi_stop_mut);*/
525
526     /* wait for thief, victim, and monitor to call this function */
527     //pthread_barrier_wait(&stop_bar);
528 }
529
530 /*
531
532     */
533 /* UTIL FUNCTIONS: */
534 /*
535
536     */
537 void rt_print_status( thread_info_t *ti)
538 {
539     printf("t_id %d status: ", ti->t_id);
540     switch( ti->status)
541     {
542         case WORK_FINISHED :
543             printf("FINISHED\n");
544             break;
545         case THIEF_SIG :
546             printf("THIEF SIG\n");
547             break;
548         case IS_THIEF :
549             printf("THIEF\n");
550             break;
551         case VICTIM_SIG :
552             printf("VICTIM SIG\n");
553             break;
554         case IS_VICTIM :
555             printf("VICTIM\n");
556             break;
557         case THREAD_WORKING :
558             printf("WORKING\n");
559     }
560 }
561
562 void rt_print_workload( thread_info_t *ti)
563 {
564     printf("t_id %d curr: %u end: %u\n",ti->t_id ,
565           ti->curr ,
566           ti->end);
567 }

```

Listing C.12: rtworksteal.h

```

1 #ifndef RTWORKSTEALH
2 #define RTWORKSTEALH
3
4 #define WORKER_COUNT 3          /* the total number of worker threads */
5
6 #define MAX_ESTIMATE 2147483647 /* FIXME not needed with limits.h */
7
8 #include <limits.h>
9 #include <time.h>
10 #include <pthread.h>
11
12 /* HAS TO BE INCLUDED LAST!!! */
13 #include "mandelbrot.h"
14
15 typedef enum { THREAD_WORKING,
16                THIEF_SIG,
17                IS_THIEF,
18                VICTIM_SIG,
19                IS_VICTIM,
20                WORK_FINISHED
21                } thread_status_t;
22
23 typedef struct thread_info_t {
24     char t_id;
25     thread_status_t status;
26     unsigned long estimated_complete;
27     unsigned int end, curr;
28
29     /* status flags */
30     char sta_t_sig, sta_v_sig, sta_stop, sta_working, sta_finished;
31
32     pthread_mutex_t stop_mut;
33     pthread_cond_t stop_cond;
34 } thread_info_t;
35
36 /*
37
38 void *rt_render_thread      ( void*);
39 void *rt_monitor_thread    ( void*);
40
41 void rt_initialise_info     ( thread_info_t *, char);
42
43 void rt_start_render_threads ();
44 void rt_join_render_threads ();
45 void rt_broadcast_finished  ();
46 void rt_distribute          ( thread_info_t *, thread_info_t *);
47 thread_info_t *rt_wait_for_thief();
48 thread_info_t *rt_find_victim ();
49
50 void rt_become_thief        ( thread_info_t *);
51 void rt_become_victim       ();
52 unsigned int rt_compute_work ( thread_info_t *);
53 void rt_update_estimate     ( thread_info_t *, unsigned long);
54
55 /*
56
57 void ws_initialise_threads  ();
58

```

```
59 void ws_start_threads      ();
60
61 /*
62     */
63 void rt_print_status( thread_info_t *);
64 void rt_print_workload( thread_info_t *);
65
66 #endif /* RTWORKSTEALH */
```